

3 Forward and Back Propagation

$$z = x_1 * w_1 + x_2 * w_2 + b$$

$$\hat{y} = a = \sigma(z)$$

$$\ell(a, y)$$

$$\frac{\partial \ell}{\partial z} = \frac{\partial \ell}{\partial a} \cdot \frac{\partial a}{\partial z}$$

$$= \left\{ -\frac{y}{a} + \frac{1-y}{1-a} \right\} * a * (1-a)$$

$$= a - y$$

$$\frac{\partial \sigma(z)}{\partial z} = \sigma(z) \cdot (1 - \sigma(z))$$

$$\frac{\partial a}{\partial z} = a \cdot (1-a)$$

$$\frac{\partial \ell}{\partial a} = \frac{y}{a} + \frac{1-y}{1-a}$$

$$z = X * W + b$$

$$\hat{y} = a = \sigma(z)$$

$$\sigma(z) = \frac{1}{1+e^{-z}}$$

$$\ell(a, y) = -[y * \log(a) + (1-y) * \log(1-a)]$$

For binary classification:

$$\ell(a, y) = -y * \log(a)$$

$$\Rightarrow \frac{\partial \ell}{\partial w_1} = x_1 \cdot \frac{\partial \ell}{\partial z} = x_1 \cdot (a-y)$$

$$\frac{\partial \ell}{\partial w_2} = x_2 \cdot \frac{\partial \ell}{\partial z} = x_2 \cdot (a-y)$$

$$\frac{\partial \ell}{\partial b} = \frac{\partial \ell}{\partial z} = (a-y)$$

$$\Rightarrow$$

$$w_1 = w_1 - \alpha * \frac{\partial \ell}{\partial w_1} = w_1 - \alpha * x_1 * (a-y)$$

$$w_2 = w_2 - \alpha * \frac{\partial \ell}{\partial w_2} = w_2 - \alpha * x_2 * (a-y)$$

$$b = b - \alpha * \frac{\partial \ell}{\partial b} = b - \alpha * (a-y)$$

Where  $\alpha$  is learning rate. The cost function is

$$J(W, b) = \frac{1}{m} * (\sum \ell(a, y))$$

Hence  $\frac{\partial J}{\partial w_1} = \frac{1}{m} * (\sum \frac{\partial \ell(a, y)}{\partial w_1})$

11/11/2025

**pra-sami**

5

## Activation Function

11/11/2025

**pra-sami**

6

## Overview

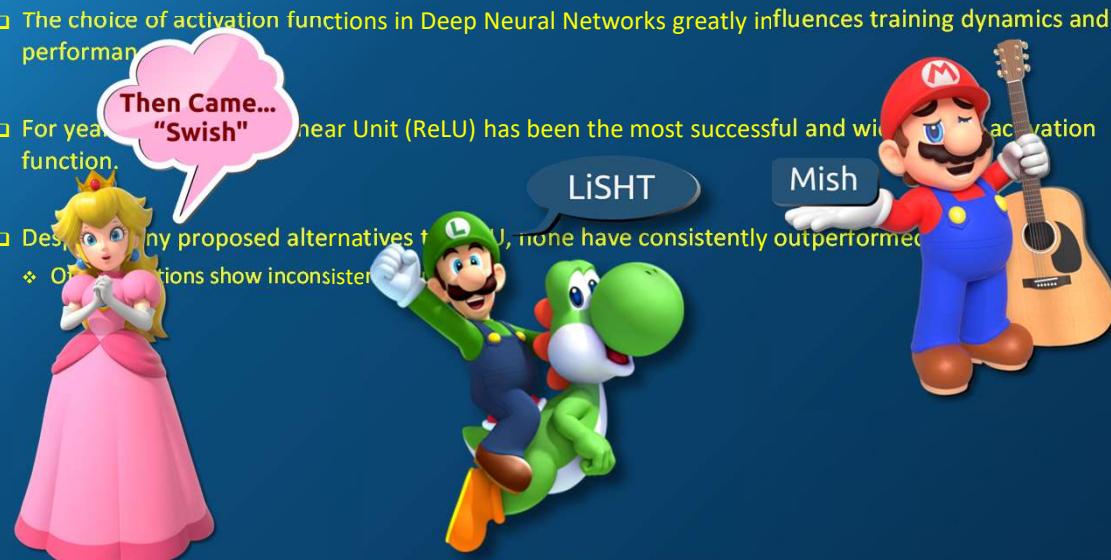
- ❑ The choice of activation functions in Deep Neural Networks greatly influences training dynamics and performance.
- ❑ For years, the Rectified Linear Unit (ReLU) has been the most successful and widely-used activation function.
- ❑ Despite many proposed alternatives to ReLU, none have consistently outperformed it:
  - ❖ Other functions show inconsistent results.

11/11/2025

7

## Overview

- ❑ The choice of activation functions in Deep Neural Networks greatly influences training dynamics and performance.
- ❑ For years, the Rectified Linear Unit (ReLU) has been the most successful and widely-used activation function.
- ❑ Despite many proposed alternatives to ReLU, none have consistently outperformed it:
  - ❖ Other functions show inconsistent results.



Linearly Scaled Hyperbolic Tangent

11/11/2025

8

## Activation Functions

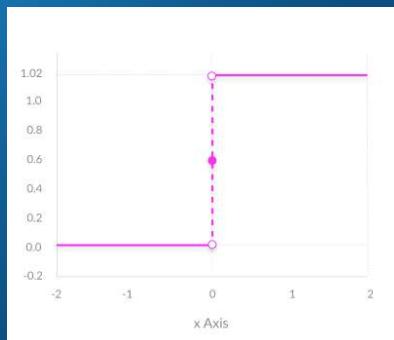
- ❑ An activation function is applied to each neuron in the network, determining whether it should be “fired” based on the relevance of its input for the model’s prediction.
- ❑ Activation functions help normalize neuron output to specific ranges:
  - ❖ Common ranges include [0, 1], [-1, 1], or other desired intervals.
- ❑ They must be computationally efficient as they are calculated for each neuron across all data instances.
- ❑ In essence, an activation function acts as a mathematical gate, turning a neuron “on” or “off.”

11/11/2025

9

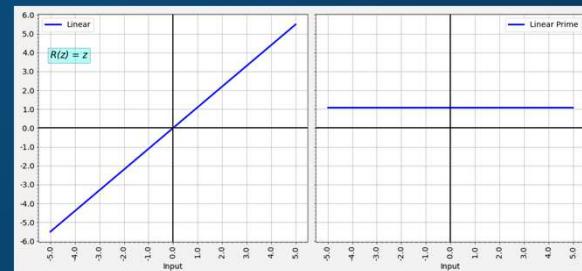
## Activation Functions

- ❑ Binary Step function



- ❑ We already seen this in previous session!!!!

- ❑ Linear Activation Function



- ❑ That will be simple linear regression!
  - ❖ There are still be some use cases...

11/11/2025

10

## Non-Linear Activation Functions

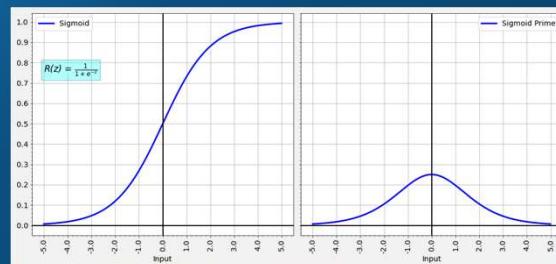
- There are many popular activation functions
  - ❖ Sigmoid / Logistic
  - ❖ Softmax
  - ❖ Tanh (Hyperbolic Tangent)
  - ❖ ReLU (Rectified Linear Unit)
  - ❖ Leaky ReLU
  - ❖ Parametric ReLU
  - ❖ ELU (Exponential Linear Unit)
  - ❖ GELU (Gaussian Error Linear Unit)
  - ❖ Swish
  - ❖ Lish
  - ❖ Mish
- Stay tuned... it's an active research area...

11/11/2025



11

## Sigmoid



- Takes a real-valued input and outputs a value between 0 and 1, i.e., [0,1].
- Easy to work with and widely used as an activation function.
- Non-linear, continuously differentiable, monotonic, with a fixed output range.
- Well-suited for binary classification tasks.

11/11/2025



12

## Sigmoid – drawbacks

- ❑ Saturation at the ends:
  - ❖ The function becomes sluggish towards either end (close to 0 or 1), causing very small gradients and making updates slow.
- ❑ Vanishing Gradient Problem:
  - ❖ Sigmoid can cause gradients to vanish during backpropagation, especially in deep networks, making it hard for the network to learn effectively.
- ❑ Non-zero-centered output:
  - ❖ This can lead to uneven gradient updates, where weights are adjusted differently for positive and negative inputs, making optimization harder.
- ❑ Gradient Saturation:
  - ❖ Near 0 or 1, it "kills" the gradients, leading to stalled learning in certain layers of the network.
- ❑ Scaling Requirement:
  - ❖ Input values often need to be scaled beforehand to avoid saturation and improve learning speed

11/11/2025



13

## Softmax Function

- ❑ Originates from physics as the Boltzmann or Gibbs distribution; formulated by physicist Ludwig Boltzmann in 1868.
- ❑ Applied to reinforcement learning by Robert Duncan Luce in 1959 in his book “Individual Choice Behavior: A Theoretical Analysis.”
- ❑ Functionality:
  - ❖ Takes a vector of N real-valued inputs and converts it into a vector of N values that sum to 1, effectively creating a probability distribution.
  - ❖ Accepts positive or negative inputs and outputs values between 0 and 1.
- ❑ Application:
  - ❖ Commonly used in the output layer of classification models to represent probabilities for each class.

11/11/2025



14

## Softmax Function

- ❑ Multi-class Logistic Regression: Softmax is essentially an extension of logistic regression for multi-class classification.
- ❑ Range: Accepts both positive and negative values as input and outputs values between 0 and 1, i.e.,  $0 \leq \text{output} \leq 1$ 
  - ❖ Converts it into a vector of N values that sum to 1, effectively representing probabilities.
- ❑ Differentiable: Softmax is differentiable everywhere, making it suitable for backpropagation in neural networks.
- ❑ Relation to Sigmoid:
  - ❖ Sigmoid is a special case of Softmax for binary classification.
- ❑ Role in Neural Networks: Softmax is typically used in the final layer of multi-layer neural networks, converting raw scores into probabilities.
- ❑ Non-linear Nature: As a non-linear function, it adds non-linearity to the network, which enhances the model's learning capability.

11/11/2025



15

## Softmax vs. Sigmoid

- ❑ Sigmoid

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

- ❑ Softmax

$$S(\vec{Z}) = \frac{e^{z_i}}{\sum_{j=1}^K e^{z_j}}$$

- ❑ For single class value will be [ 0, x], Softmax

$$S(\vec{Z}) = \frac{e^{z_i}}{\sum_{j=1}^K e^{z_j}}$$

$$S(\vec{Z}) = \frac{e^{z_1}}{e^{z_1} + e^{z_2}}$$

$$S(X) = \frac{e^x}{e^0 + e^x}$$

$$S(X) = \frac{e^x}{1 + e^x}$$

$$S(X) = \frac{1}{1 + e^{-x}}$$

Like Sigmoid Activation function, Vanishing Gradient is still a problem!

11/11/2025



16

## Softmax vs. Argmax

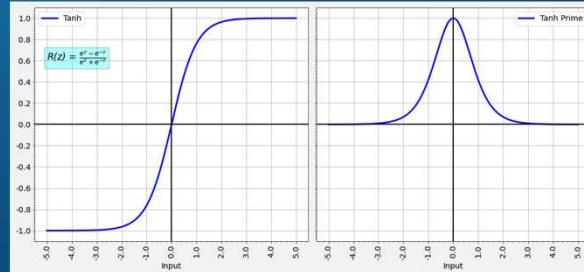
- ❑ Both work the same way, Softmax is expected to be a differentiable alternative to argmax
- ❑ Argmax returns index of highest value and no idea about other values.
- ❑ It is common to train using the Softmax

11/11/2025



18

## Tanh



- ❑ A shifted version of the Sigmoid function, mapping inputs to a range of  $[-1,1]$
- ❑ Non-linear and zero-centered: This zero-centered nature helps center data around zero, making it easier for the next layer to learn.
- ❑ Stronger Gradient: Tanh has a stronger gradient than Sigmoid, with steeper derivatives, making learning somewhat faster.
- ❑ Useful for Hidden Layers: Commonly used in hidden layers to help stabilize the network's activations.

11/11/2025



19

## Tanh

❑ Advantages of the Tanh Function:

- ❖ Handles Negative and Zero Inputs Well: Maps negative inputs to negative values and zero inputs near zero, aiding in better data representation.
- ❖ Differentiable and Monotonic: The function is differentiable, and while it's monotonic, its derivative is not, which can aid in diverse learning behaviors.
- ❖ Faster Convergence:
  - Steeper gradients than the Sigmoid function enable faster learning.
  - Zero-centered output helps balance the gradients, making optimization easier.

❑ Disadvantages of the Tanh Function:

- ❖ Vanishing Gradient Problem: Tanh still suffers from the vanishing gradient issue at extreme values, which can slow down learning, especially in deep networks.
- ❖ Mixed Views in Research: Some studies suggest Tanh is better than Sigmoid, while others find limited or conditional advantages. The debate over its effectiveness continues.

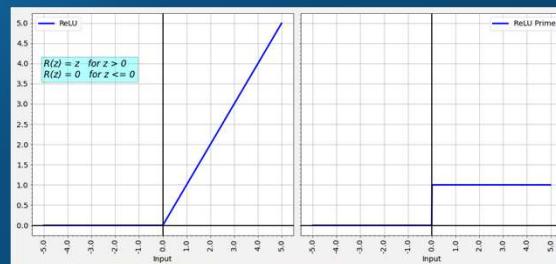
❑ Practical Use: Tanh can be a good starting point in the early design stages for hidden layers in neural networks.

11/11/2025



20

## Rectified Linear Units (ReLU)



- ❑ Non-linear function (almost)
- ❑ Better performance than Sigmoid or Tan in almost all models
- ❑ It avoids and rectifies vanishing gradient problem.
- ❑ ReLU is less computationally expensive than tanh and sigmoid because it involves simpler mathematical operations.
- ❑ Suitable for Hidden layers only.

11/11/2025



21

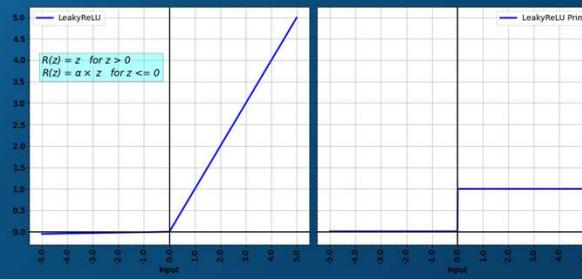
## Rectified Linear Units (ReLU)

- Some gradients can be fragile during training and can die.
- For activations in the region ( $x < 0$ ) of ReLU , gradient will be zero
  - ❖ Weights will not get adjusted during descent
  - ❖ Neurons which go into that state will stop responding to variations in error/ input
  - ❖ Dying ReLU problem
- The range of ReLU is  $[0, \infty]$ 
  - ❖ Can blow up the activation

11/11/2025

22

## Leaky ReLU



- a variant of the ReLU (Rectified Linear Unit) activation function
  - ❖ Attempt to fix the “dying ReLU” problem
- Leaky ReLU allows a small, non-zero slope (usually a small constant like 0.01) for negative values.
  - ❖  $R(z_i) = \begin{cases} z_i & \text{if } z_i \geq 0 \\ a_i \cdot z_i & \text{if } z_i < 0 \end{cases}$

11/11/2025

23

## Leaky ReLU

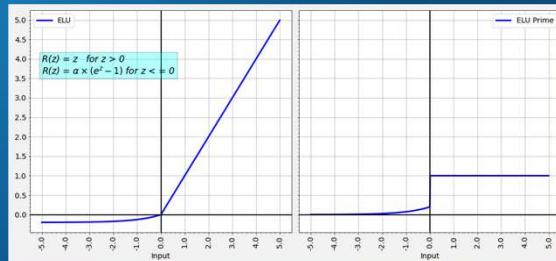
- Prevents Dying Neurons
- Improves Training Efficiency
- Simplicity and Effectiveness
- Disadvantages:
  - ❖ Small Negative Gradient
  - ❖ Hyperparameter Tuning
  - ❖ Not Zero-Centered
  - ❖ Potential for Exploding Gradients
  - ❖ Limited Improvement over ReLU
- These limitations are generally minor and can often be mitigated with proper tuning or alternative activation functions, depending on the specific neural network architecture

11/11/2025



24

## Exponential Linear Unit (ELU)



$$F(z_i) = \begin{cases} z_i & \text{if } z_i \geq 0 \\ \alpha * (e^{z_i} - 1) & \text{if } z_i < 0 \end{cases}$$

$\alpha$  is generally 1.0 but can be tuned.

- Converges faster ; Has alpha constant which should be positive number
- ELU is a strong alternative to ReLU.
  - ❖ Provides a smooth curve for negative values
- Unlike to ReLU, ELU is zero-centered, as it outputs negative values for negative inputs and positive values for positive inputs, helping improve gradient flow
- For  $x > 0$ , it can blow up the activation with the output range of  $[0, \infty]$ .

11/11/2025



25

## Exponential Linear Unit (ELU)

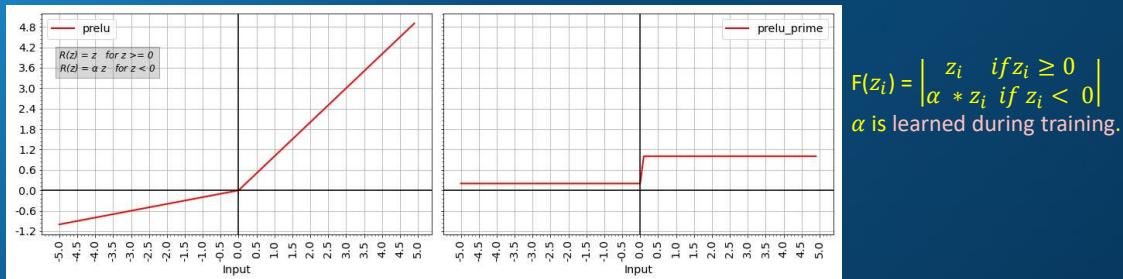
Advantages	Disadvantages
<ul style="list-style-type: none"> <li>□ Prevents Dying Neurons           <ul style="list-style-type: none"> <li>❖ Allows a small negative output (instead of zero)</li> </ul> </li> </ul>	<ul style="list-style-type: none"> <li>□ Computational Cost           <ul style="list-style-type: none"> <li>❖ Requires computing an exponential function</li> </ul> </li> </ul>
<ul style="list-style-type: none"> <li>□ Zero-Centered           <ul style="list-style-type: none"> <li>❖ Can lead to faster convergence</li> </ul> </li> </ul>	<ul style="list-style-type: none"> <li>□ Potential for Exploding Gradients           <ul style="list-style-type: none"> <li>❖ Uses an exponential function for negative inputs. For very large negative inputs, the function can produce large gradients</li> </ul> </li> </ul>
<ul style="list-style-type: none"> <li>□ Smoothness           <ul style="list-style-type: none"> <li>❖ ELU is continuous and differentiable everywhere</li> </ul> </li> </ul>	<ul style="list-style-type: none"> <li>□ Hyperparameter Tuning</li> </ul>
<ul style="list-style-type: none"> <li>□ Better Representation for Negative Values           <ul style="list-style-type: none"> <li>❖ Network can learn better representations for data with negative correlations.</li> </ul> </li> </ul>	<ul style="list-style-type: none"> <li>□ Slower Convergence           <ul style="list-style-type: none"> <li>❖ It may not always converge as quickly as other activation functions, like Swish or SELU</li> </ul> </li> </ul>
	<ul style="list-style-type: none"> <li>□ Limited Performance in Certain Architectures</li> </ul>

11/11/2025



26

## Parameterized ReLU



- A Parametric Rectified Linear Unit, or PReLU, is an activation function that generalizes the traditional rectified unit with a slope for negative values.
- The intuition is that different layers may require different types of nonlinearity.
- $\alpha$  is initialized and then updated during training as part of learning process
- For complex networks this may lead to overfitting.

11/11/2025



27

## Parameterized ReLU

- ❑ Pick your own parameter ( $\alpha$ ). Layers learn the parameter.
- ❑ In experiments with convolutional neural networks, PReLus for the initial layer have more positive slopes, i.e. closer to linear.
  - ❖ Since the filters of the upper layers are edge or texture detectors,
  - ❖ This shows a circumstance where positive and negative responses of filters are respected.
- ❑ In contrast, deeper layers have smaller coefficients
  - ❖ Model becomes more discriminative at later layers
  - ❖ While it wants to retain more information at earlier layers.

11/11/2025



28

## Challenges with ReLU

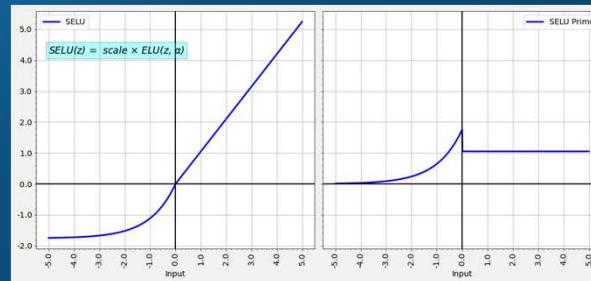
- ❑ The consistent problem is that its derivative is 0 for half of the values of the input  $x$  in the Function, i.e.  $f(x)=\max(0,x)$
- ❑ As parameter update algorithm, could used Stochastic Gradient Descent and other optimizers
  - ❖ If the parameter itself is 0, then that parameter will never be updated as it just assigns the parameter back to itself
  - ❖ Leading close to 40% Dead Neurons in the Neural network environment where  $z$  is negative
  - ❖ Various substitutes like Leaky ReLU, Parameterized ReLU, have unsuccessfully tried to devoid it of this issue.

11/11/2025



29

## Scaled ELU (SELU)



- Activation was introduced in a 2017 paper by Klambauer et al
- Properly initialized, the networks will self-normalize
  - ❖ Each layer's output will roughly be zero-centered with standard deviation equal to one
- Helps prevent the vanishing or exploding gradients problems

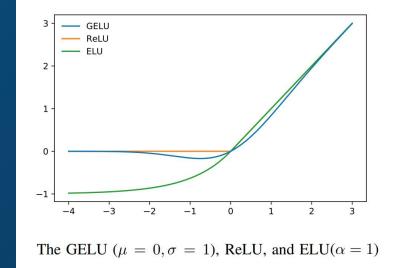
11/11/2025



30

## Gaussian Error Linear Unit (GELU)

- A smooth, continuous activation function
- Popular in deep learning, especially in transformer-based models like BERT and GPT.
- Combines the benefits of both the ReLU function and probabilistic interpretations based on the Gaussian distribution,
  - ❖ Providing a smoother approach to activation.
- Contrary to the ReLU, GELU weights its inputs by their value instead of thresholding them by their sign
- The GELU activation function is  $x * \Phi(x)$ ,
  - ❖ where  $\Phi(x)$  : the standard Gaussian cumulative distribution function refer scipy's norm.cdf(x)
  - ❖  $GELU(x) = x P(X \leq x) = x \Phi(x) \approx 0.5 x [1 + \tanh\{\sqrt{2/\pi} (x + 0.04475 x^3)\}]$
  - ❖ Or  $x \sigma(1.702 x)$

The GELU ( $\mu = 0, \sigma = 1$ ), ReLU, and ELU( $\alpha = 1$ )

11/11/2025



31

## Gaussian Error Linear Unit (GELU)

- ❑ Smoother nonlinearity
  - ❖ Combination of the gaussian distribution and the error function provides a smoother curve for both positive and negative values
- ❑ Probabilistic interpretation
  - ❖ Behaves like a soft version of relu
- ❑ Zero-centered
- ❑ Smooth handling of negative values
  - ❖ Probabilistic nature means that negative values are not entirely squashed to zero (like in relu) but instead undergo a smooth transformation

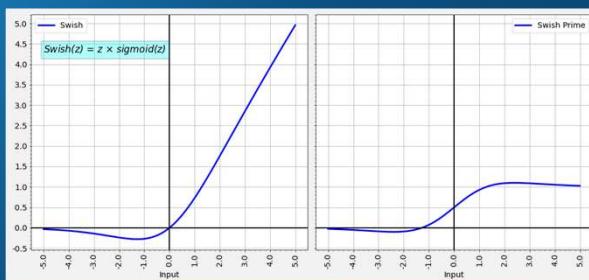
- ❑ Advantages:
  - ❖ Improved gradient flow
  - ❖ Zero-centered output
  - ❖ Better performance in complex architectures
- ❑ Disadvantages
  - ❖ Computational complexity
  - ❖ Slower training : due to its smooth, non-linear nature and additional computational overhead, GELU may lead to slower training times compared to relu and its variants
  - ❖ Difficult hyperparameter tuning
  - ❖ Limited adoption
  - ❖ Possible exploding gradients

11/11/2025



32

## Swish



- ❑ Google Brain Team proposed a new activation function:
  - ❖  $f(x) = x \cdot \text{sigmoid}(x)$
- ❑ Experiments show that Swish tends to work better than ReLU on deeper models across a number of challenging data sets
  - ❖ Simply replacing ReLUs with Swish units improves top-1 classification accuracy on ImageNet by 0.9% for Mobile NASNetA and 0.6% for Inception-ResNet-v2
- ❑ The simplicity of Swish and its similarity to ReLU make it easy for practitioners to replace ReLUs with Swish units in any neural network.
- ❑ Swish is a smooth, non-monotonic function that consistently matches or outperforms ReLU on deep networks

11/11/2025



33

## Swish

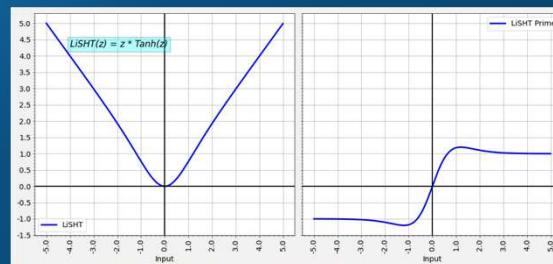
- ❑ Unbounded above and bounded below
  - ❖ Non-monotonic attribute that actually creates the difference
- ❑ We can train deeper Swish networks than ReLU networks when using BatchNorm (Ioffe & Szegedy, 2015) despite having gradient squishing property
- ❑ With MNIST data set, when Swish and ReLU are compared, both activation functions achieve similar performances up to 40 layers.
- ❑ Swish outperforms ReLU by a large margin in the range between 40 and 50 layers
  - ❖ For less than 40 layers, performance is comparable
- ❑ In very deep networks, Swish achieves higher test accuracy than ReLU.
- ❑ Swish outperforms ReLU on every batch size, suggesting that the performance difference between the two activation functions remains even when varying the batch size.
- ❑ Gradient descent problem was still there may be to a lesser degree!

11/11/2025



34

## LiSHT Activation Function



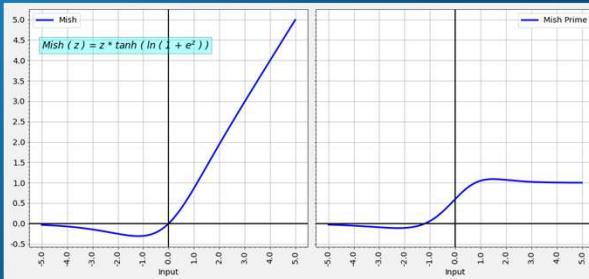
- ❑ The function scale the non-linear Hyperbolic Tangent ( Tanh ) function by a linear function
  - ❖ Help tackle the dying gradient problem
- ❑ According to paper it has outperformed Swish on a number of problems

11/11/2025



35

## Mish



$$\begin{aligned} f(z) &= z * \tanh(\text{softplus}(z)) \\ &= z * \tanh(\ln(1 + e^z)) \end{aligned}$$

- Inspired by Swish and has been shown to outperform it in a variety of computer vision tasks
- Mish was “found by systematic analysis and experimentation over the characteristics that made Swish so effective”.
- Mish seems to be the best activation in stock,
  - ❖ But jury is still out

11/11/2025



36

## Reflect...

- Which of the following is a common activation function used in last layer of deep neural networks?
  - ❖ A) Linear activation
  - ❖ B) Step function
  - ❖ C) Sigmoid function
  - ❖ D) Exponential function
- Answer: C) Sigmoid function
- What is the vanishing gradient problem in deep neural networks?
  - ❖ A) The problem of too many layers in the network
  - ❖ B) The problem of exploding gradients during training
  - ❖ C) The problem of slow convergence during training
  - ❖ D) The problem of very negligible gradients in early layers
- Answer: D) The problem of very negligible gradients in early layers
- What is the purpose of the softmax activation function in the output layer of a classification neural network?
  - ❖ A) To introduce non-linearity
  - ❖ B) To convert logits into probabilities
  - ❖ C) To prevent overfitting
  - ❖ D) To reduce the dimensionality of the output
- Answer: B) To convert logits into probabilities
- In deep learning, what does the term "epoch" refer to during training?
  - ❖ A) A complete pass through the training dataset
  - ❖ B) The number of layers in the neural network
  - ❖ C) The learning rate of the optimizer
  - ❖ D) The size of the mini-batch used for training
- Answer: A) A complete pass through the training dataset

11/11/2025



37

## Reflect...

- What is the role of activation functions in deep neural networks?
  - ❖ A) To normalize the input data
  - ❖ B) To compute the loss function
  - ❖ C) To introduce non-linearity into the network
  - ❖ D) To reduce overfitting
- Answer: C) To introduce non-linearity into the network
- What does the term "backpropagation" refer to in the context of neural networks?
  - ❖ A) The process of adjusting the weights of the network based on the prediction error
  - ❖ B) The process of training the network using labeled data
  - ❖ C) The process of selecting the optimal hyperparameters for the network
  - ❖ D) The process of initializing the weights of the network
- Answer: A) The process of adjusting the weights of the network based on the prediction error

11/11/2025



38

## Next Session...



Loss / Cost Optimization

Stochastic Gradient Descent (SGD) and others

Momentum Learning Rates

Adaptive Learning Rates

11/11/2025



39



11/11/2025

**pra-sami**