



## Vanishing and Exploding Gradients Gradient Check

Deep Neural Networks  
Session 11  
Pramod Sharma  
[pramod.sharma@prasami.com](mailto:pramod.sharma@prasami.com)

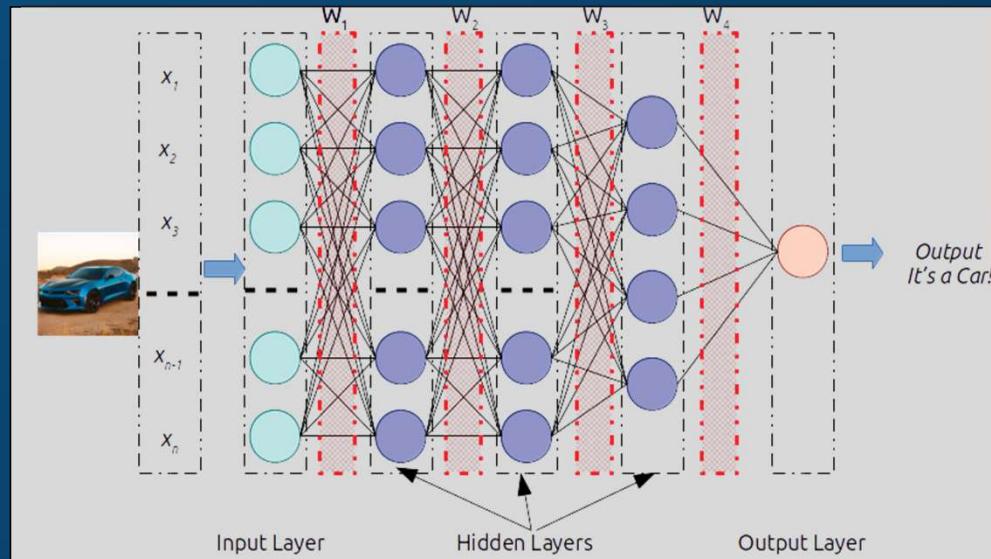
Neural Networks face problems of vanishing or exploding gradients

Deeper the network more are the chances of the gradients becoming smaller and smaller or keep growing...

## Weights multiplied...

- We know that :
  - ❖  $z = X * W + b$
  - ❖  $\hat{y} = a = \sigma(z)$
  - ❖  $a_1 = \sigma(a_0 \cdot W_1)$

- That in multilayer network
  - ❖  $\hat{y} = \sigma(\sigma(\sigma(\sigma(a_0 \cdot W_1) \cdot W_2) \cdot W_3) \cdot W_4) \dots$
  - ❖ For explanation purpose assume  $\sigma(z) = z$  (say ReLU)
  - ❖  $y = W_1 \cdot W_2 \cdot W_3 \cdot W_4$
  - ❖ so any change in  $y$  will result in  $W_1 * W_2 * W_3 * W_4$  times  $y$  in layer 1
  - ❖ Longer the chain, more  $W_s$  will be multiplied.



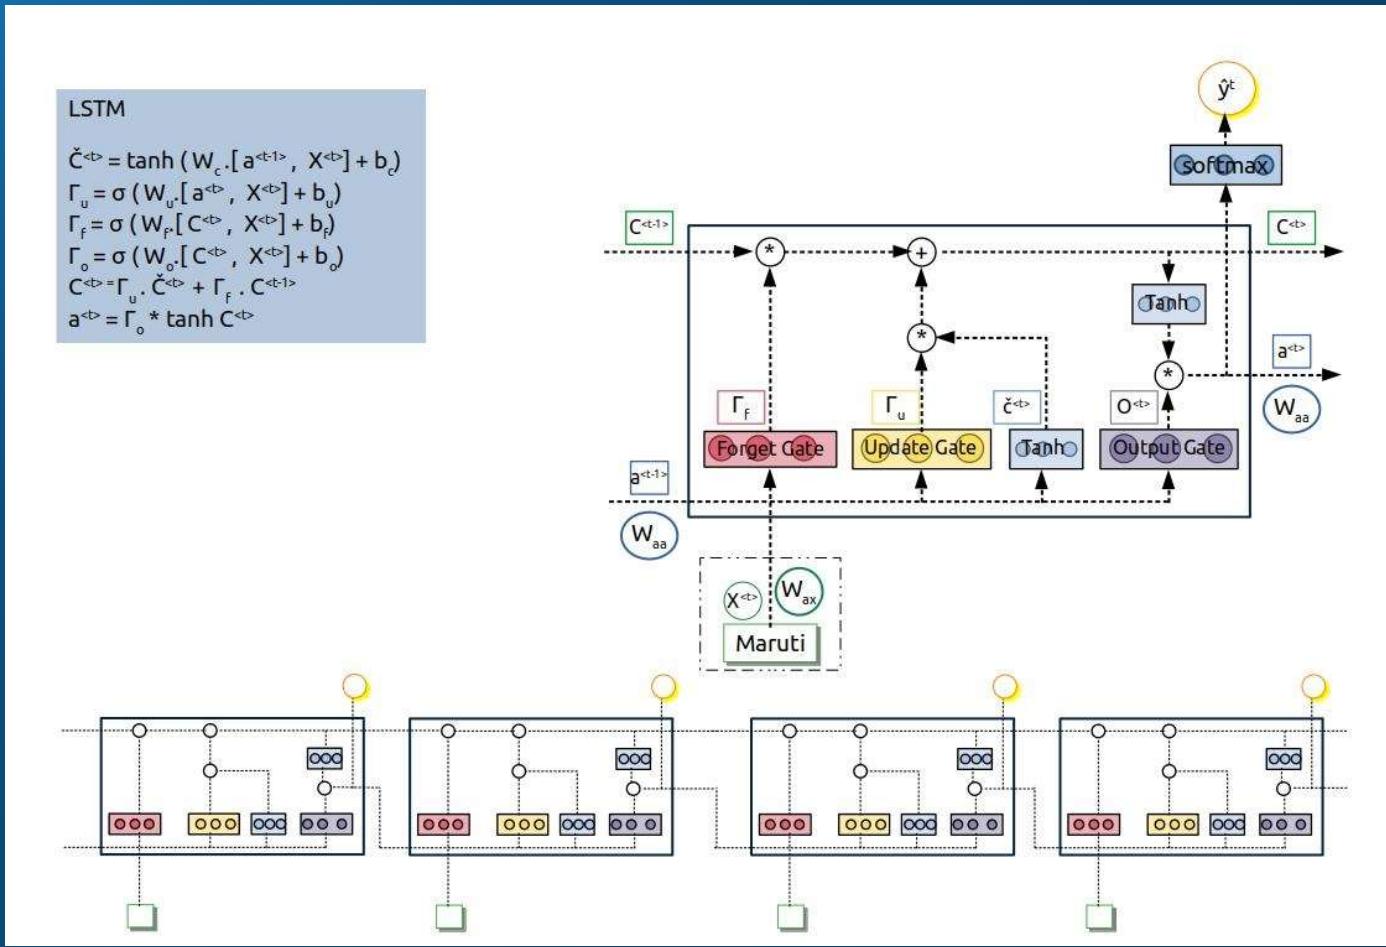
## More Layers... More problems

- Assume we have 150 layers
- Also assume our weight is say 1.1  
 $=> 1.1^{150} = 1.6 \text{ million}$
- On the other hand assume our weight is 0.9  
 $=> 0.9^{150} = 1.4 e^{-7}$

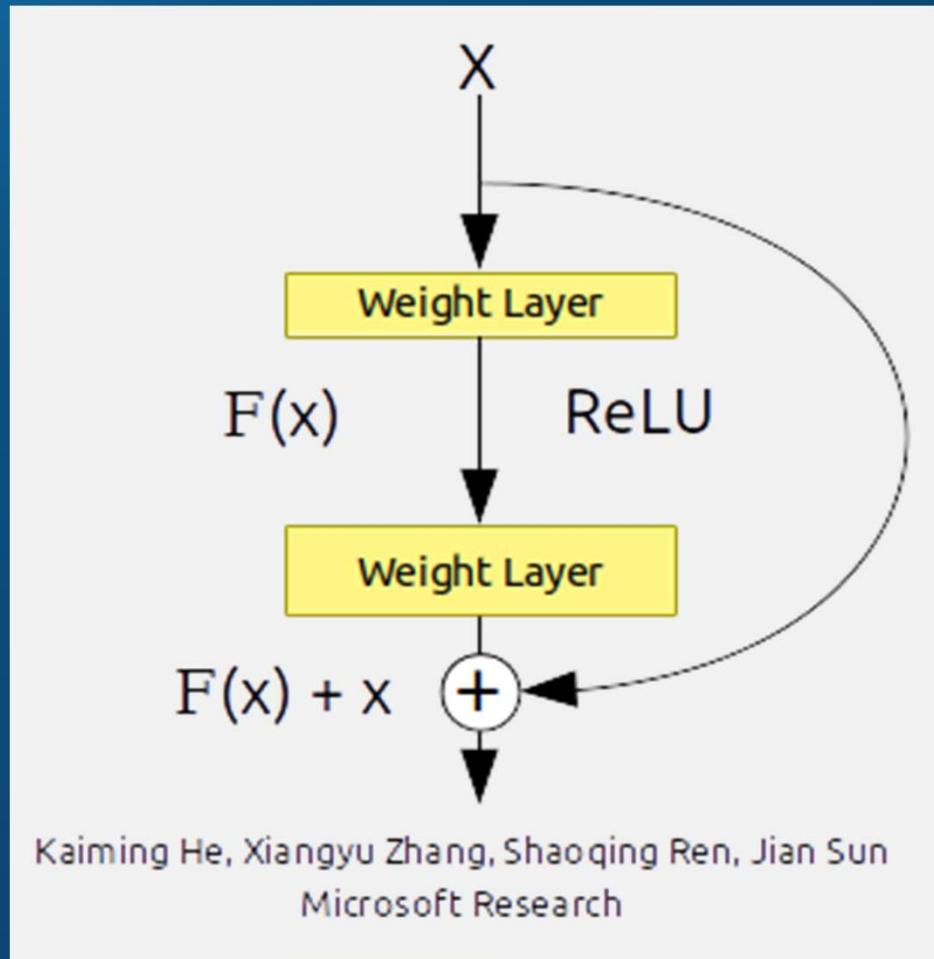
## It's a Severe Problem...

- No silver bullet solution....
- There is multi-prong approach to it...
- First, Initialise your weights as close to 1 as possible ( not 1)
- It is found that for tanh activation function
  - ❖ Divide by  $\sqrt{\text{number of nodes in the previous layer}}$
- Some cases :  $\frac{2}{\sqrt{\text{number of nodes in the previous layer}}}$
- In ReLU ,  $\frac{2}{\sqrt{\text{number of nodes in the previous layer} + \text{number of nodes in current layer}}}$
- Some literature suggest, even  $\frac{K}{\sqrt{\text{number of nodes in the previous layer}}}$ ;  $K$  is a another parameter to tune

## Change Your Architecture - LSTM



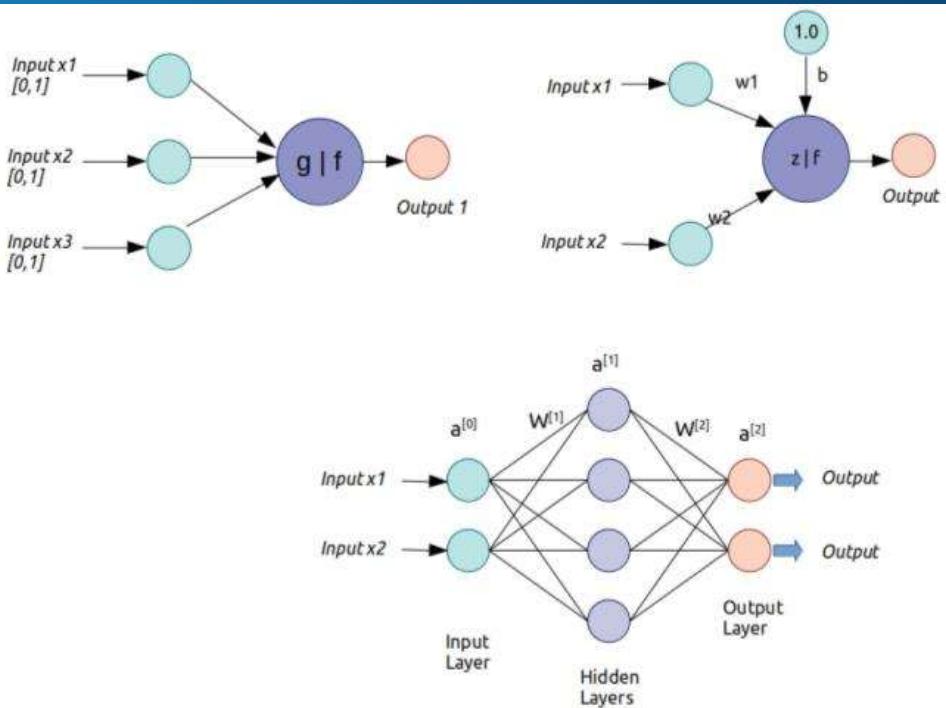
## As in ResNet



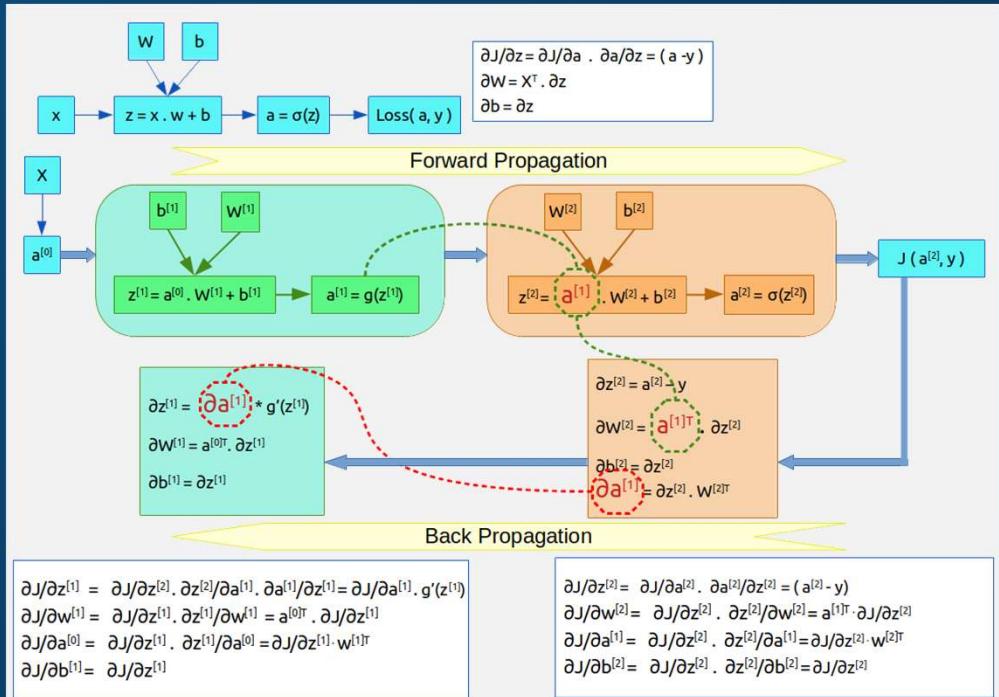
## Gradient Check

# Gradient Checks

## □ Recap



## □ The math



## Loss Function and its Derivative

- The loss for our prediction  $\hat{y}$  with respect to the true labels  $y$  is given by:

$$L(\hat{y}, y) = -y \cdot \log \hat{y} - (1-y) \cdot \log (1-\hat{y})$$

- For all samples:

$$J(\hat{y}, y) = -\frac{1}{m} \sum_{i \in m} y_i \cdot \log \hat{y}_i - (1-y_i) \cdot \log (1-\hat{y}_i)$$

Where  $\hat{y} = \sigma(a \cdot W + b)$

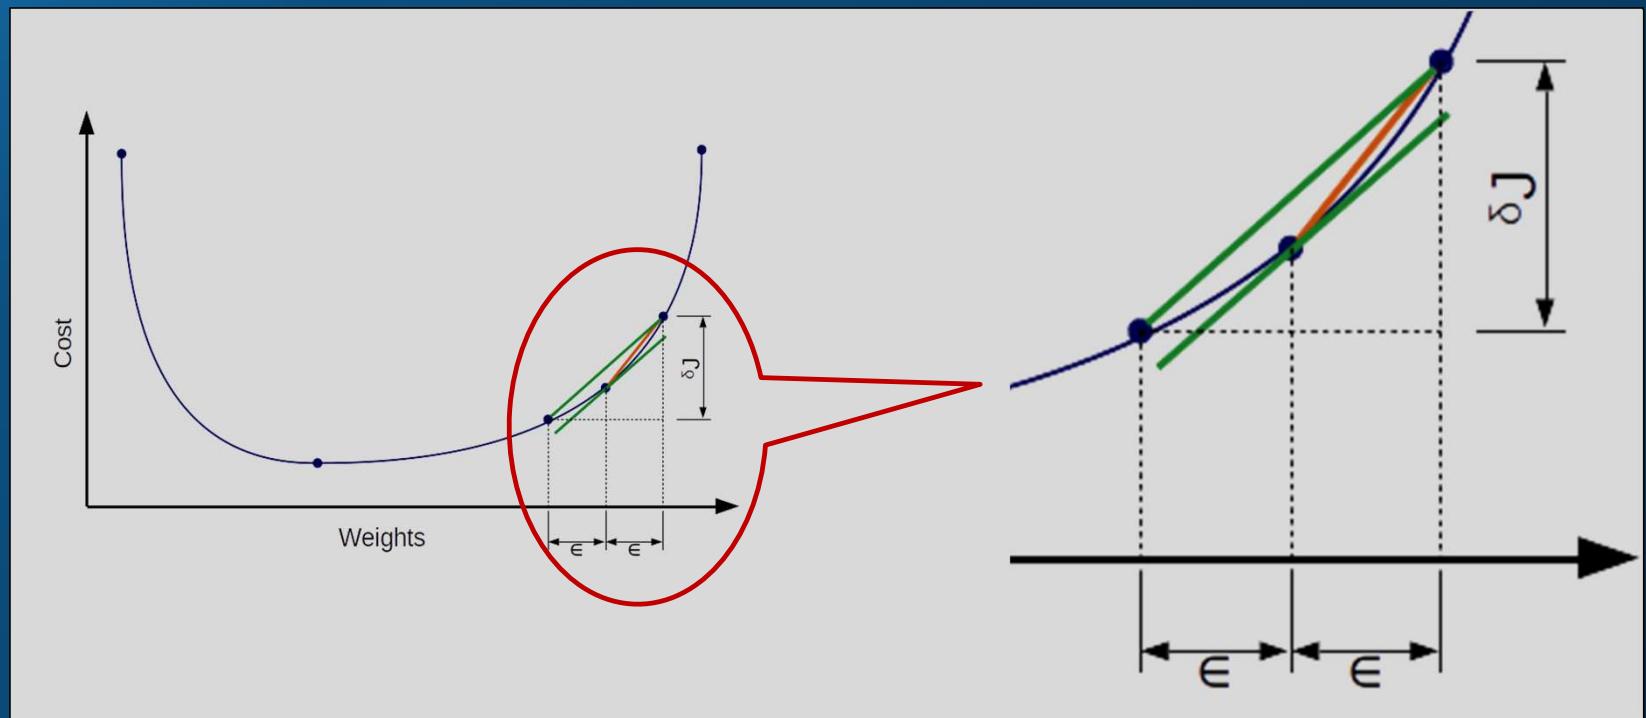
- Therefore, we can say that:

$$J(\hat{y}, y) = J(W, b) = J(W_1, W_2, W_3, \dots, W_n, b_1, b_2, b_3, \dots, b_n)$$

## Calculation of derivative

### □ Use the centered formula

- ❖ The formula you may have seen for the finite difference approximation when evaluating the numerical gradient is not as good as centered formula



## Gradient Checking

- ❑ Also called “Grad Check”
- ❑ Do it to verify the model’s math(Debug) only

❖ Too heavy for training, switch off once the model is verified.

- ❑ For all values of Ws and Bs, we can calculate:

$$\delta\theta_{approx} = J(W_1, \dots, W_i + \varepsilon, \dots, W_n, b_1, b_2, b_3, \dots, b_n) - J(W_1, \dots, W_i - \varepsilon, \dots, W_n, b_1, b_2, b_3, \dots, b_n) / (2 * \varepsilon)$$

- ❑ To check if  $\delta\theta_{approx}$  and  $\delta\theta$  are close

$$\frac{\|\delta\theta_{approx} - \delta\theta\|_2}{\|\delta\theta_{approx}\|_2 + \|\delta\theta\|_2}$$
 is very small

- ❑ For  $\varepsilon = 1e-7$ 
  - ❖ Relative error  $> 1e-2$  usually means the gradient is probably wrong
  - ❖  $1e-2 >$  relative error  $> 1e-4$  should make you feel uncomfortable
  - ❖  $1e-4 >$  relative error is usually okay for objectives with kinks. But if there are no kinks (e.g. use of tanh nonlinearities and softmax), then  $1e-4$  is too high.
  - ❖  $1e-7$  and less you should be happy

## Grad Check Steps

- Recall: Our model has all weights and biases stored
  - ❖ Model = { "W<sub>1</sub>": ..., "b<sub>1</sub>": ..., "W<sub>2</sub>": ..., "b<sub>2</sub>": ..., ... ..., "W<sub>n</sub>": ..., "b<sub>n</sub>": ... }
  - ❖ We have implemented our forward prop and back prop
- Step 1 : Pick model and convert all weights and biases into a vector  $\theta$
- Step 2: Similarly pick  $\delta\theta$  and convert to a vector  $\delta\theta$
- Step 3: for each of the value in the vector  $\theta$ 
  - ❖ Make copy of  $\theta$  and  $\delta\theta$
  - ❖ Increase  $\theta_i$  to  $\theta_i + \varepsilon$
  - ❖ Calculate  $J_+$  ( Cost with increased  $\theta$ )
  - ❖ Similarly calculate  $J_-$  ( Cost with decreased  $\theta$ )
  - ❖ Use  $J_+$  and  $J_-$  to calculate  $\delta\theta_{approx}$
  - ❖ Calculate  $\delta\theta$  as usual
  - ❖ Find error

## Remember to Turn off Dropout/Augmentations

- When performing gradient check, remember to turn off any non-deterministic effects in the network, such as dropout, random data augmentations, etc.
- Otherwise these can clearly introduce huge errors when estimating the numerical gradient
- Downside → you wouldn't be gradient checking them
  - ❖ It might be that dropout isn't back-propagated correctly)

## Use Double Precision

- ❑ A common pitfall is using single precision floating point to compute gradient check. It is often the case that you might get high relative errors (as high as  $1e-2$ ) even with a correct gradient implementation
- ❑ There are cases when relative errors plummet from  $1e-2$  to  $1e-8$  by switching to double precision

## Stick Around Active Range of Floating Point

- What Every Computer Scientist Should Know About Floating-Point Arithmetic  
[https://docs.oracle.com/cd/E19957-01/806-3568/ncg\\_goldberg.html](https://docs.oracle.com/cd/E19957-01/806-3568/ncg_goldberg.html)
- In neural nets it is common to normalize the loss function over the batch
- If gradients per datapoint are very small, then *additionally* dividing them by the number of data points starts to give very small numbers
- Always print the raw numerical/analytic gradient,
  - ❖ Make sure that the numbers you are comparing are not extremely small (e.g. roughly 1e-10 and smaller in absolute value)
  - ❖ If so, you may want to temporarily scale your loss function up by a constant to bring them to a “nicer” range where floats are more dense - ideally on the order of 1.0, where your float exponent is 0.

## Kinks in the Objective Function

- One source of inaccuracy to be aware of during gradient checking is the problem of *kinks*
- Kinks refer to non-differentiable parts of an objective function
  - ❖ Such as ReLU  $\max(0, x)$ , or the SVM loss, Maxout neurons, etc.
- Consider gradient checking the ReLU function at  $x = -1e-6$
- Since  $x < 0$ , the analytic gradient at this point is exactly zero.
- However, the numerical gradient would suddenly compute a non-zero gradient because  $f(x + h)$  might cross over the kink (e.g. if  $h > 1e-6$ ) and introduce a non-zero contribution.
- You might think that this is a pathological case, but in fact this case can be very common.
- Keeping track of the identities of all “winners” in a function of form  $\max(x, y)$ ; That is, was  $x$  or  $y$  higher during the forward pass. If the identity of at least one winner changes when evaluating  $f(x + h)$  and then  $f(x - h)$ , then a kink was crossed and the numerical gradient will not be exact

## Grad Check

- Phew... too lengthy calculations....
- Good News!
  - ❖ Both **Tensorflow** and **Torch** have **autograd** implementation for us. So in real implementation we will be using those functions for gradient checking
- Caution!
  - ❖ This check is resource hungry
  - ❖ Once the verification is done, comment/switch off the code
- Deeper the network → the higher the relative errors
  - ❖ For the input data for a 10-layer network, a relative error of 1e-2 might be okay because the errors build up on the way
  - ❖ Conversely, an error of 1e-2 for a single differentiable function likely indicates incorrect gradient

Next Session...

The process of learning the parameters...

Finding good hyper-parameters...

