




Batch Normalization

Deep Neural Network
Session 13
Pramod Sharma
pramod.sharma@prasami.com


Batch Normalization

- It definitely helps to normalize input data
- Gradient converges faster



Batch Normalization

- What about hidden layer?
- After all activations from previous layer are inputs for current layer...



- Will it help if we normalize the hidden layers too?

Batch Normalization

- Batch normalization (also known as batch norm) [by Sergey Loffe and Christian Szegedy in 2015]
 - ❖ Make artificial neural networks faster
 - ❖ More stable through normalization of the input layer by re-centering and re-scaling
- Wider choices of hyper-parameter...
 - ❖ Higher Learning Rates: Can use much larger learning rates without training diverging
 - ❖ Less Sensitive to Weight Initialization: Networks are more robust to poor initializations
 - ❖ Reduced Need for Regularization: The inherent noise in batch statistics acts as a regularizer
 - ❖ Faster Convergence: All of the above leads to needing fewer epochs to get good results
- In theory, it's normalizing activation values of the respective layers
- In practice, it works better if we normalize 'z'
 - ❖ Look at the documentation for details

Batch Normalization

- In General, any Z^i can be normalized

$$\text{mean } \mu = \frac{\sum Z^i}{m}$$

$$\text{std } \sigma^2 = \frac{1}{m} \sum (Z^i - \mu)^2$$

- $Z^i \text{ Norm} = \frac{Z^i - \mu}{\sqrt{\sigma^2}}$
- $\hat{z} = \gamma \cdot Z^i \text{ Norm} + \beta$
- where γ and β are parameters we can train
- if $\gamma = \frac{1}{\sqrt{\sigma^2}}$ and $\beta = \frac{\mu}{\sqrt{\sigma^2}}$; $Z^i \text{ Norm} = \hat{z}$

Batch Normalization

- In General, any Z^i can be normalized

$$\text{mean } \mu = \frac{\sum Z^i}{m}$$

$$\text{std } \sigma^2 = \frac{1}{m} \sum (Z^i - \mu)^2$$

$$z^i_{\text{Norm}} = \frac{z^i - \mu}{\sqrt{\sigma^2}}$$

$$\hat{z} = \gamma \cdot z^i_{\text{Norm}} + \beta$$

- where γ and β are parameters, we can Train

- if $\gamma = \frac{1}{\sqrt{\sigma^2}}$ and $\beta = \frac{\mu}{\sqrt{\sigma^2}}$; $Z^i_{\text{Norm}} = \hat{z}$

Instead of using z^i_{Norm} , researchers realized that its better to derive \hat{z} with two trainable parameters.

Intuition is that by normalizing z , we are introducing bias in the system. Hence it makes sense to train these parameters

Batch Normalization

- In General, any Z^i can be normalized
- mean $\mu = \frac{\sum Z^i}{m}$
- std $\sigma^2 = \frac{1}{m} \sum (Z^i - \mu)^2$
 - ❖ $Z^i \text{ Norm} = \frac{Z^i - \mu}{\sqrt{\sigma^2 + \varepsilon}}$
 - ❖ $\hat{z} = \gamma \cdot Z^i \text{ Norm} + \beta$
- where γ and β are parameters, we can train
- if $\gamma = \sqrt{\sigma^2 + \varepsilon}$ and $\beta = \mu$; $Z^i \text{ Norm} = \hat{z}$


Lets add a small ε to prevent zero divide error...

Batch Normalization


□ Notes:

- ❖ Batch norm is used along with mini batches
- ❖ Batch norm is applied to the batch under consideration only irrespective of other mini batches

□ Where does it fit in overall scheme?





Batch Normalization



Batch Normalization


For $a^{[2]}$ all $a^{[1]}$ are acting as input features






Batch Normalization

- Forward and back propagation with batch norm:




Our standard equation to calculate z_1 .




Batch Normalization

- Forward and back propagation with batch norm:




Calculate \hat{z}_1 , based on β_1, γ_1




Batch Normalization

- Forward and back propagation with batch norm:




Apply activation function $g_1(\hat{z}_1)$




Batch Normalization

- Forward and back propagation with batch norm:





Calculate z_2 as usual



Batch Normalization


- Forward and back propagation with batch norm:



We know how to calculate \hat{z}_2

Batch Normalization


- Forward and back propagation with batch norm:



We also know how to calculate a_2

Batch Normalization


- Forward and back propagation with batch norm:



$$\begin{aligned}\beta_1 &= \beta_1 - \alpha \cdot \partial \beta_1 \\ \gamma_1 &= \gamma_1 - \alpha \cdot \partial \gamma_1\end{aligned}$$


Using the gradient descent,
update β 's, γ 's along with
W's and b's

$$\begin{aligned}\beta_2 &= \beta_2 - \alpha \cdot \partial \beta_2 \\ \gamma_2 &= \gamma_2 - \alpha \cdot \partial \gamma_2\end{aligned}$$



Batch Normalization

- Forward and back propagation with batch norm:




$$\begin{aligned}\beta_1 &= \beta_1 - \alpha \cdot \partial \beta_1 \\ \gamma_1 &= \gamma_1 - \alpha \cdot \partial \gamma_1\end{aligned}$$

$$\begin{aligned}\beta_2 &= \beta_2 - \alpha \cdot \partial \beta_2 \\ \gamma_2 &= \gamma_2 - \alpha \cdot \partial \gamma_2\end{aligned}$$

One more thing, since we are normalizing our Z 's, keeping
 b 's in the equation does not make any sense now.
 Being the constant it will get eliminated!!

Batch Normalization

- Forward and back propagation with batch norm:




$$\begin{aligned}\beta_1 &= \beta_1 - \alpha \cdot \partial \beta_1 \\ \gamma_1 &= \gamma_1 - \alpha \cdot \partial \gamma_1\end{aligned}$$

$$\begin{aligned}\beta_2 &= \beta_2 - \alpha \cdot \partial \beta_2 \\ \gamma_2 &= \gamma_2 - \alpha \cdot \partial \gamma_2\end{aligned}$$

And at test/validation time using a exponentially weighted average!


So while training do not forget to save exponentially weighted values or simply running average!!



Batch Normalization

- Forward and back propagation with batch norm:

Too many calculation steps...



Don't worry...

Most frameworks have one line code to do it.

Batch Normalization – Code Sample

```
model = tf.keras.models.Sequential(  
    [  
        tf.keras.layers.RNN( keras.layers.LSTMCell(units), input_shape=(None, input_dim) ),  
        tf.keras.layers.BatchNormalization(),  
        tf.keras.layers.Dense(output_size),  
    ]  
)  
  
class Net(nn.Module):  
    def __init__(self):  
        super(Net, self).__init__()  
        self.dense1 = nn.Linear(in_features=320, out_features=50)  
        self.dense1_bn = nn.BatchNorm1d(50)  
        self.dense2 = nn.Linear(50, 10)
```

- And it is applied to mini batches **only....**
- Batch Norm can be updated using any of the optimization functions...

Batch Normalization



Remember β, γ are parameters you train!

At test time, we use these fixed running_mean and running_var values for normalization.

Reflect...

- What is the primary purpose of Batch Normalization in deep learning?
 - ❖ A) To prevent overfitting
 - ❖ B) To reduce the number of parameters in the model
 - ❖ C) To accelerate training and reduce internal covariate shift
 - ❖ D) To increase the depth of the neural network
- Answer: C) To accelerate training and reduce internal covariate shift
- At which stage is Batch Normalization applied in a neural network?
 - ❖ A) After the input layer
 - ❖ B) After the activation function
 - ❖ C) Before the loss calculation
 - ❖ D) Before or after the activation function, depending on the implementation
- Answer: D) Before or after the activation function, depending on the implementation
- Which of the following is a key step in Batch Normalization?
 - ❖ A) Normalizing the gradient updates
 - ❖ B) Normalizing the activations by subtracting the batch mean and dividing by the batch standard deviation
 - ❖ C) Initializing weights to zero
 - ❖ D) Adding noise to the input data
- Answer: B) Normalizing the activations by subtracting the batch mean and dividing by the batch standard deviation
- What are the two learnable parameters introduced in Batch Normalization?
 - ❖ A) Gamma and Beta
 - ❖ B) Alpha and Beta
 - ❖ C) Theta and Gamma
 - ❖ D) Sigma and Mu
- Answer: A) Gamma and Beta

Next Session... Recurrent Neural Networks