



Gates, GRU

Deep Neural Network
Session 19
Pramod Sharma
pramod.sharma@prasami.com

Agenda



Where we are

Vanishing and exploding gradients

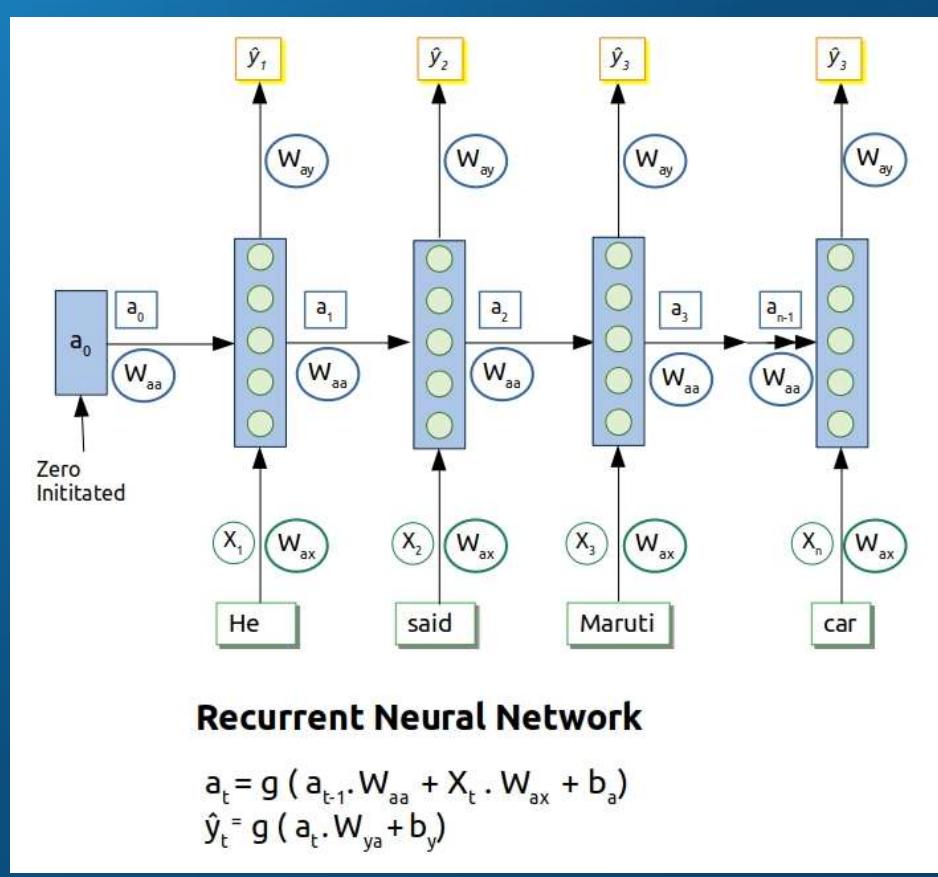
Keeping Things Stable

GRU Units

Recurrent Neural Networks (RNNs)

- Recurrent Neural Networks take the previous output or hidden states as inputs
- The composite input at time “t” has some historical information about the happenings at time $T < t$.
- RNNs are useful as their intermediate values (state) can store information about past inputs for a time that is not fixed a priori
- Note that the weights are shared over time
- Essentially, copies of the RNN cell are made over time (unrolling/ unfolding), with different inputs at different time steps

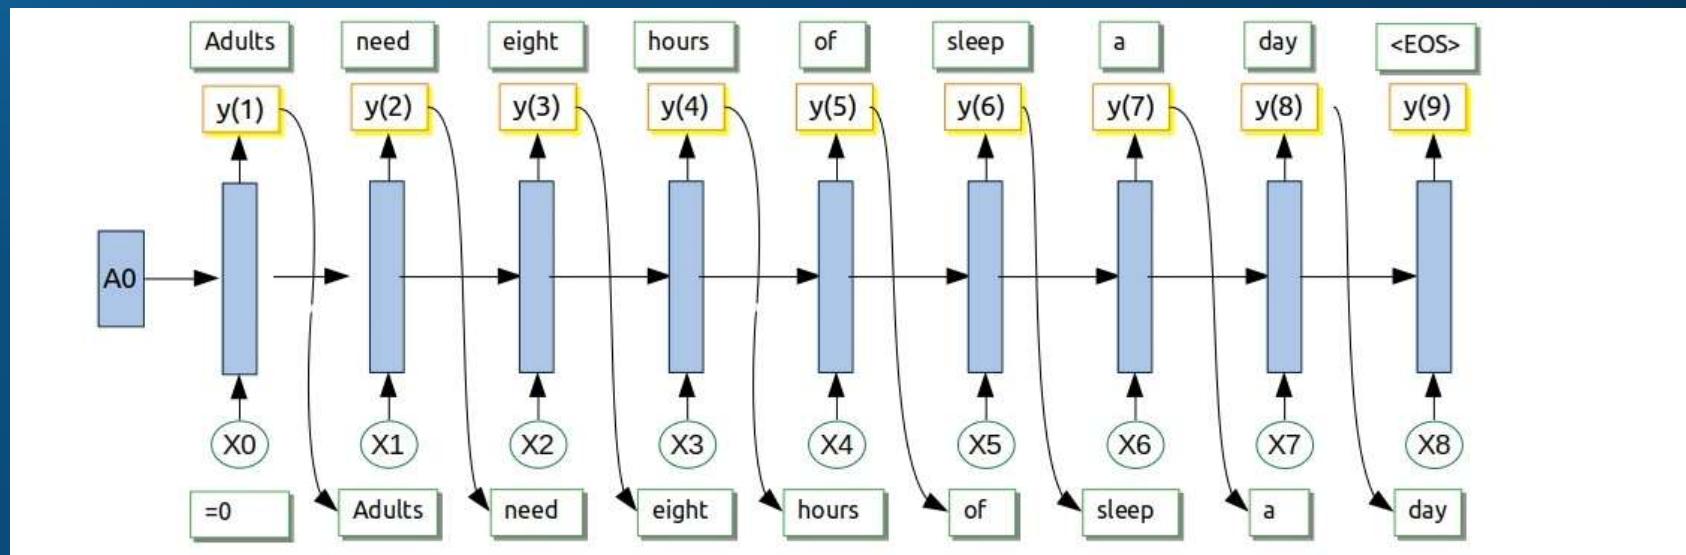
That's is Recurrent Neural Network...



□ We worked out math too....

RNN Model – Sampling from Trained Model

- And the Cost functions
- $J(\hat{y}, y) = \sum \ell(\hat{y}, y)$ BTW: these cost functions are also known as Jacobians
- $J(\hat{y}, y) = -\frac{1}{m} \sum y * \log(\hat{y})$
 - ❖ Which we will be minimizing



Gradient - a Difficult Terrain

- Have seen how to compute the gradient descent update for using backprop
- In case of RNN the backprop is through time
- At times, gradient descent completely fails because either they explode or vanish
- It's hard to learn dependencies over long time windows
- How to learn long-term dependencies?

“

Sentences can be tricky...

“As he crossed toward the pharmacy at the corner he involuntarily turned his head because of a burst of light that had ricocheted from his temple, and saw, with that quick smile with which we greet a rainbow or a rose, a blindingly white parallelogram of sky being unloaded from the van—a dresser with mirrors across which, as across a cinema screen, passed a flawlessly clear reflection of boughs sliding and swaying not arboreally, but with a human vacillation, produced by the nature of those who were carrying this sky, these boughs, this gliding façade.”

”

Vladimir Nabokov, “The Gift.” 96 words sentence...

Vanishing and Exploding Gradients

Vanishing Gradient / Exploding Gradient

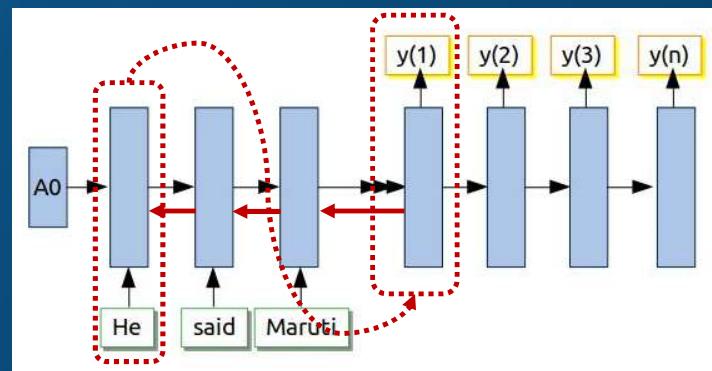
- What happens to the magnitude of the gradients as we back propagate through many layers?
 - ❖ If the weights are small, the gradients shrink exponentially
 - ❖ If the weights are big the gradients grow exponentially
- Typical feed-forward neural nets can cope with these exponential effects because they only have a few hidden layers
- We can manage gradients by initializing the weights very carefully in feed-forward networks
 - ❖ We have already experienced by using appropriate
- Is it applicable to RNNs as well?

Vanishing Gradient / Exploding Gradient

- In an RNN trained on long sequences (e.g. 100 time steps) the gradients can easily explode or vanish.
- Even with good initial weights, it's very hard to detect that the current target output depends on an input from many time-steps ago
 - ❖ So RNNs have difficulty dealing with long-range dependencies

Why Gradients Explode or Vanish

- Recall the RNN for machine translation
 - ❖ For example, we read an entire English sentence, and then has to output its French translation

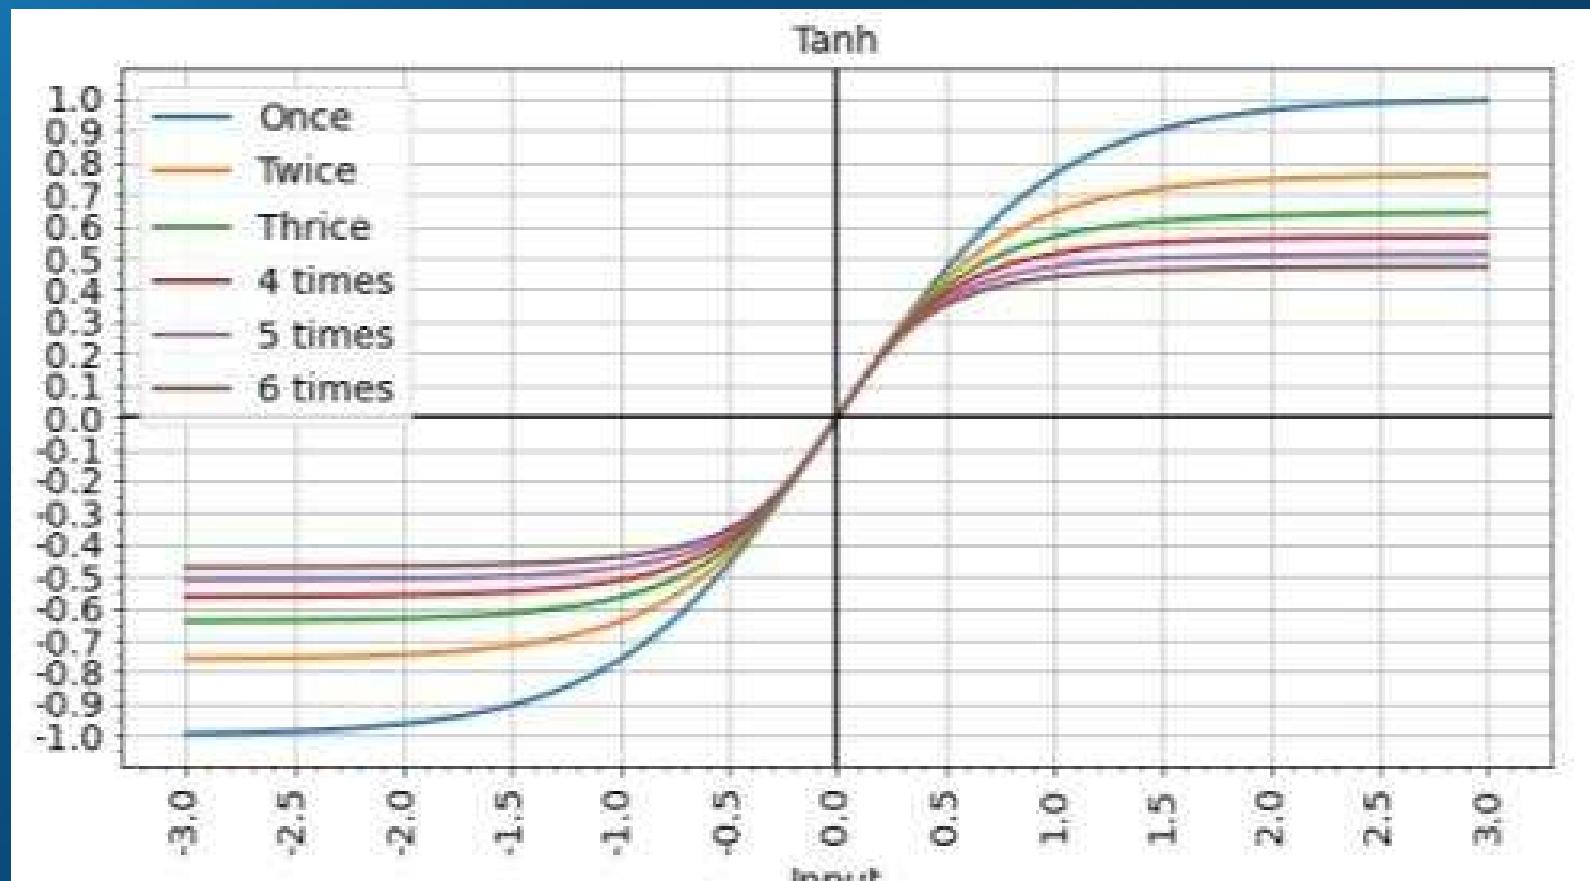


- A typical sentence length is 20 words. This means there's a gap of 20 time steps between when we see some information and when we need it.
- The derivatives need to travel over this entire pathway

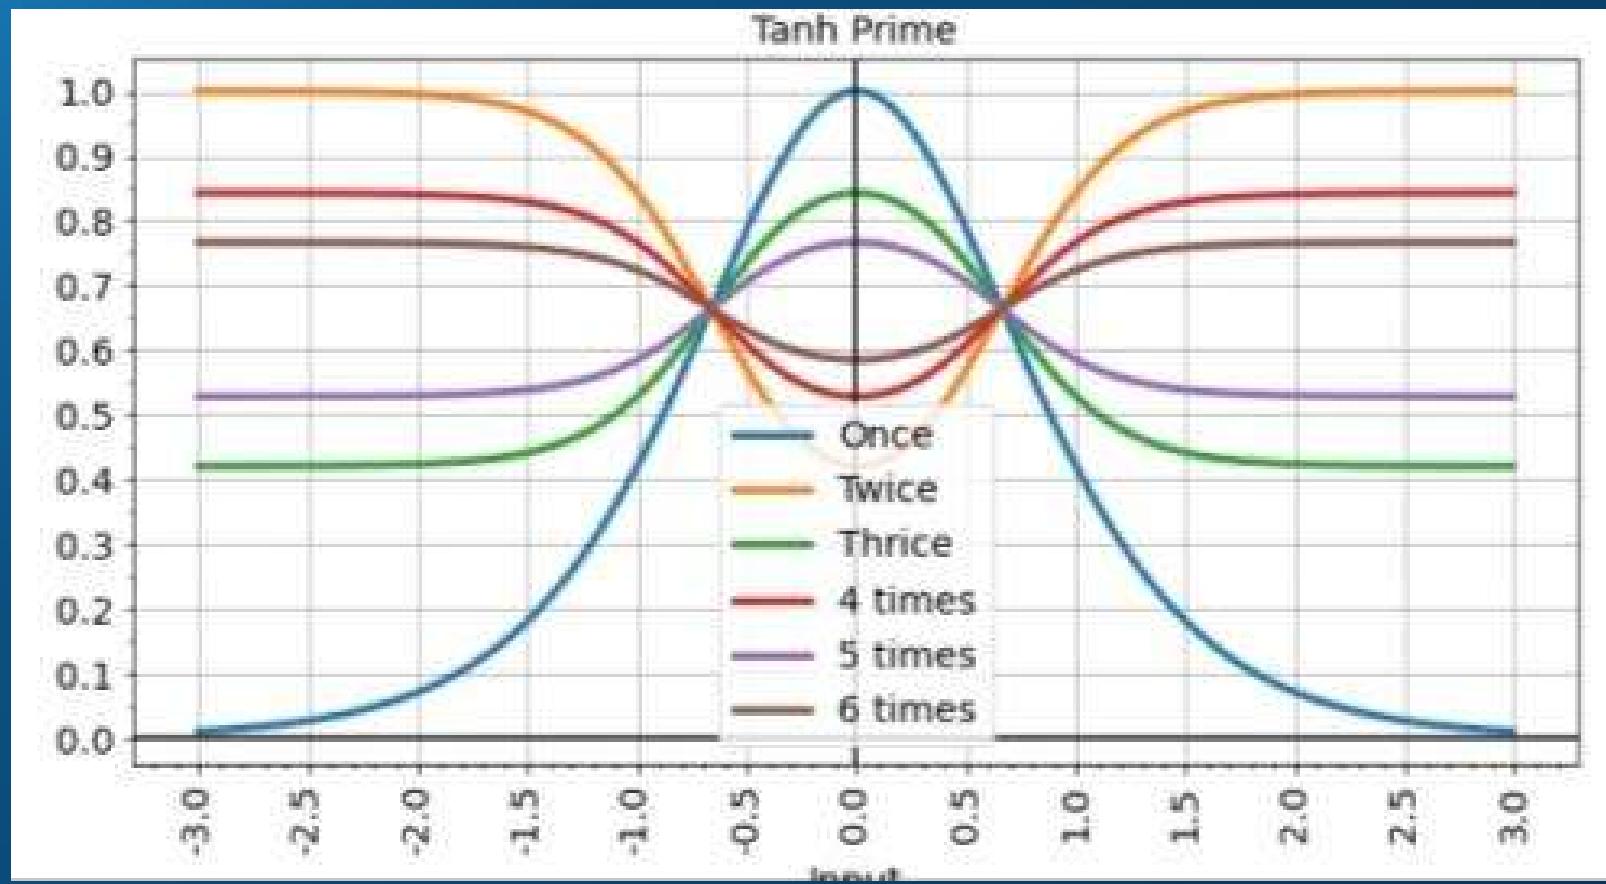
Why Gradients Explode or Vanish...

- Please recall:
 - ❖ $z = X \cdot W + b$
 - ❖ $\hat{y} = a = \sigma(z)$
 - ❖ $a_1 = \sigma(a_0 \cdot W_1)$
- That through the time steps will be
 - ❖ $\hat{y} = \sigma(\sigma(\sigma(\sigma(a_0 \cdot W_1) \cdot W_2) \cdot W_3) \cdot W_4)$
- In backprop, we will be carrying $L(\hat{y}, y)$ through the activation function iteratively...
- Longer the chain... more iterations on the Ws...

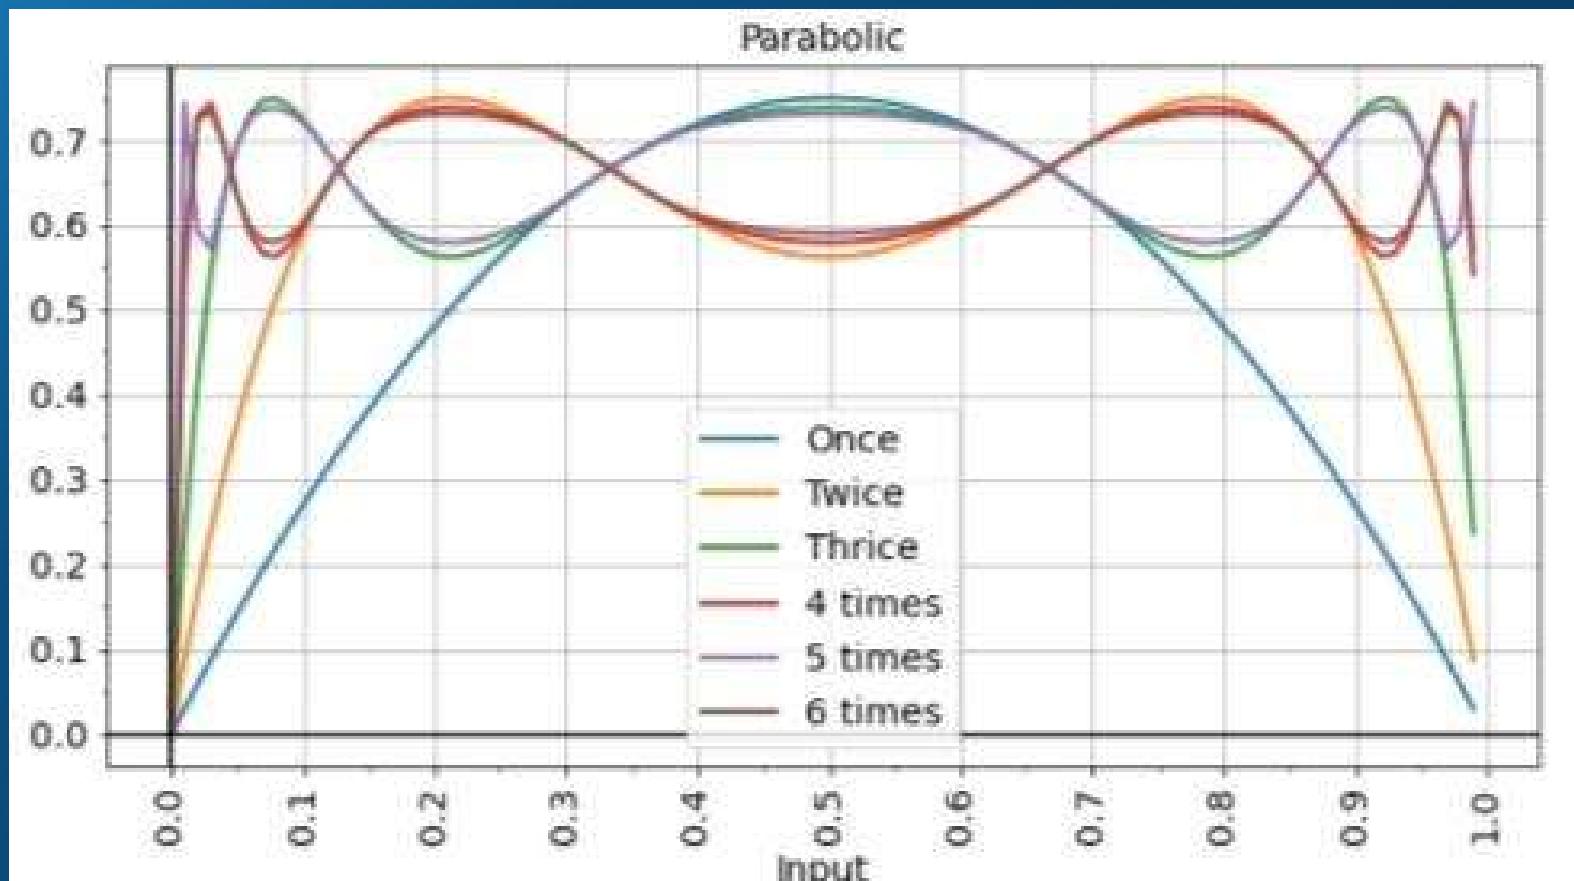
Iterated functions...



Iterated functions...

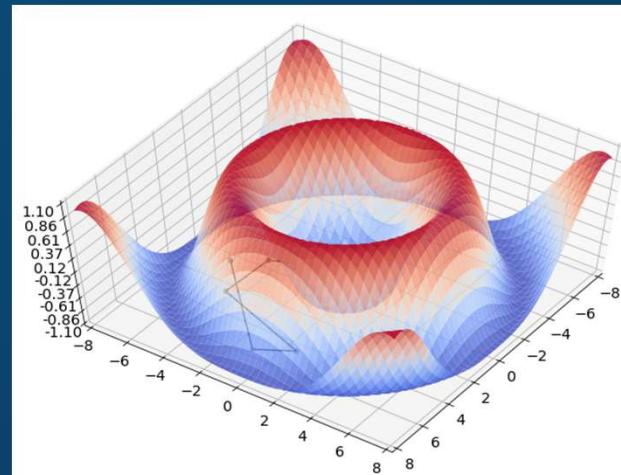
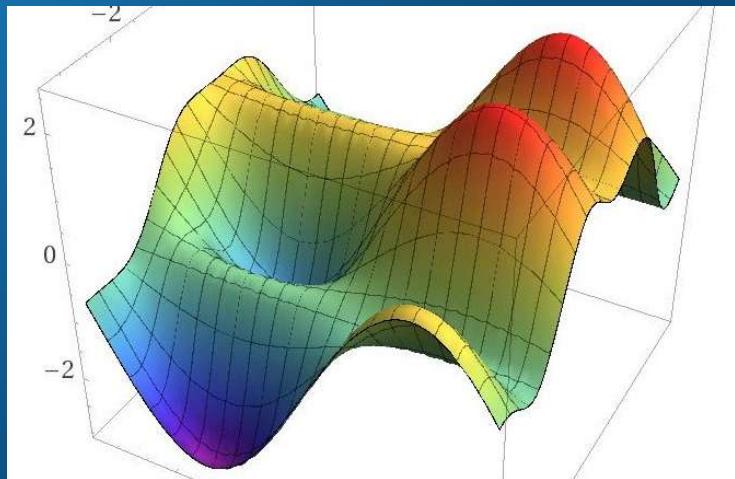


Iterated functions...



Why Gradients Explode or Vanish

- Let's imagine an RNN's behavior as a dynamical system, which has various slopes and valleys



- Within one of the colored regions, the gradients vanish because even if you move a little, you still wind up at the same valley
- If you're on the boundary, the gradient blows up because moving slightly moves you from one side to another and subsequently to other valley

Keeping Things Stable

Keeping Things Stable - Gradient Clipping

- Clip the gradient
- Clip the gradient 'g' so that it has a norm of at most 'η':
 - ❖ if $\|g\| > \eta$: then $g = \eta * \frac{g}{\|g\|}$
 - ❖ Where 'η' is another parameter you may want to tune
- The gradients are biased, but at least they don't blow up.

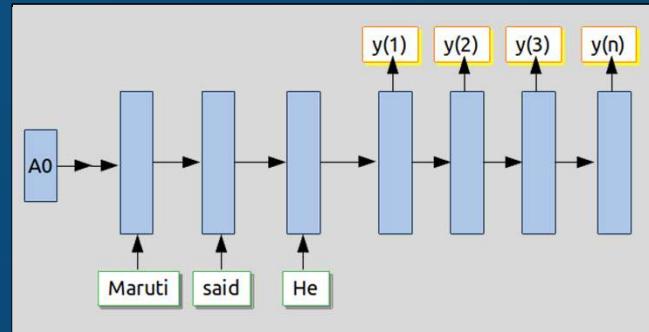
Keeping Things Stable - Reverse the Input Sequence

- Applicable in similar languages
 - ❖ Hindi → Marathi,
 - ❖ English → French,
 - ❖ Spanish → Portuguese

- No point in using situation like
 - ❖ Hindi → Mandarin or ‘Hiragana’ or ‘Kanji’ ; may be ‘Katakana’

- This way, there’s only one time step between the first word of the input and the first word of the output.

- The network can first learn short-term dependencies between early words in the sentence, and then long-term dependencies between later words.



Keeping Things Stable – Identity initialization

- Redesign the architecture, since the exploding/vanishing problem highlights a conceptual problem with vanilla RNNs
- The hidden units are a kind of memory. Therefore, their default behavior should be to keep their previous value.
 - ❖ The function at each time step should be close to the identity function.
 - ❖ It's hard to implement the identity function if the activation function is nonlinear!
- If the function is close to the identity, the gradient computations are stable

Keeping Things Stable – Identity initialization

- The identity RNN architecture : [Le et al., 2015. A simple way to initialize recurrent networks of rectified linear units.]
 - ❖ The activation functions are all ReLU,
 - ❖ Recurrent weights are initialized to the identity matrix
- Proof: This simple initialization trick achieved some neat results;
- For instance, it was able to classify MNIST digits which were fed to the network one pixel at a time, as a length-784 sequence

Applicability

- Discussed three mechanism for training RNNs
 - ❖ All pretty widely used.
 - ❖ But the identity initialization trick actually eludes to something much more fundamental
 - ❖ Keep their previous value, unless it is necessary to change
- Ask: the ability to preserve information over time until it's needed

GRU Units

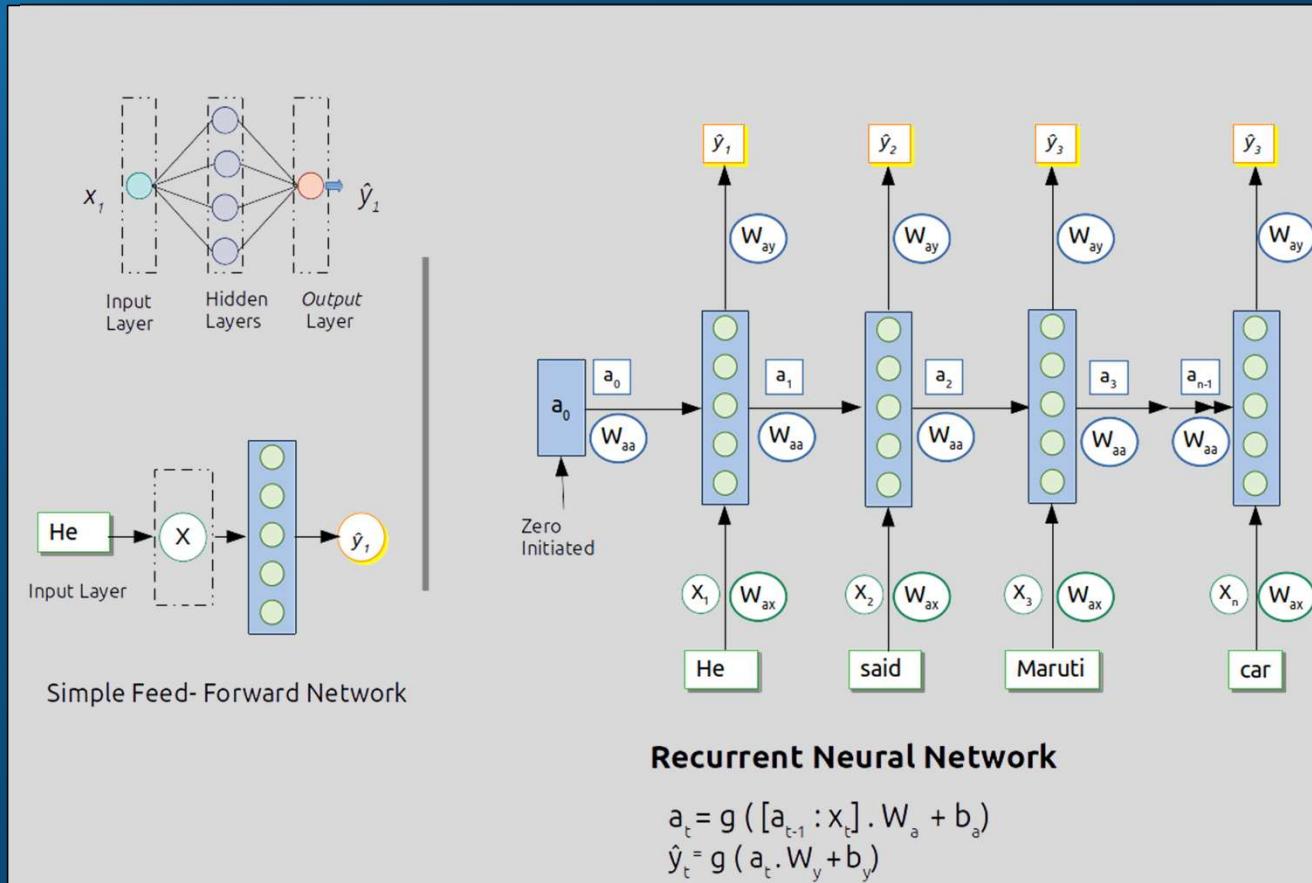
Long Short Term Memory (LSTM)

- Architecture designed to make it easy to remember
- Keep it until it's needed.
- The name refers to the idea:
 - ❖ The activations of a network correspond to short-term memory,
 - (Changing very fast with every new incoming record)
 - ❖ The weights correspond to long-term memory.
- If the activations can preserve information over multiple time steps
 - ❖ That makes them long-term short-term memory
- It's composed of memory cells which have controllers governing when to store or forget information

Gated Recurrent Unit

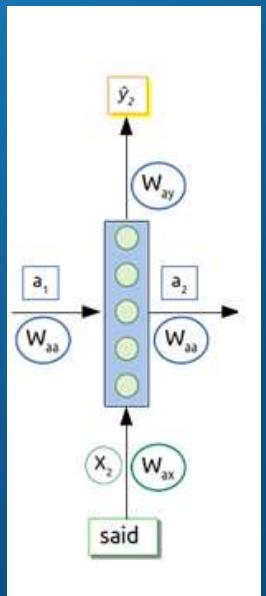
- Before we get to LSTM, lets look at its simplified version...
- Introduced by Cho, et al. in 2014, and Chung et al. in 2014 in their respective papers
- GRU (Gated Recurrent Unit)

Gated Recurrent Unit



Converted our simple feed forward network to Recurrent Network

Gated Recurrent Unit



□ RNN Equation

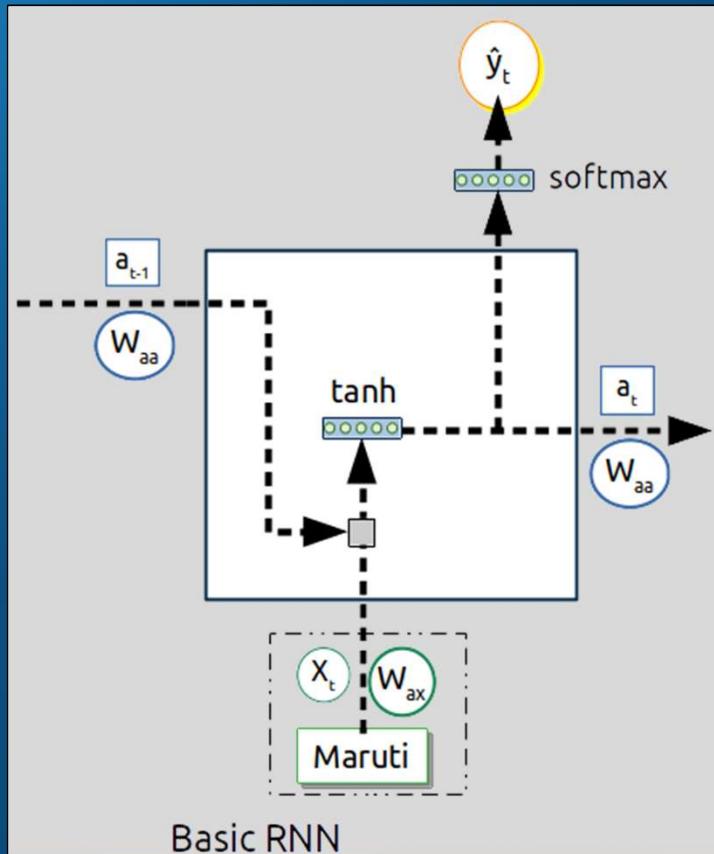
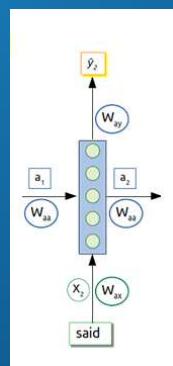
- ❖ $a_t = g_1 ([a_{t-1} : x_t] \cdot W_a + b_a)$

and

- ❖ $\hat{y}_t = g_2 (a_t \cdot W_y + b_y)$

□ We look at an alternate way to paint the network

RNN

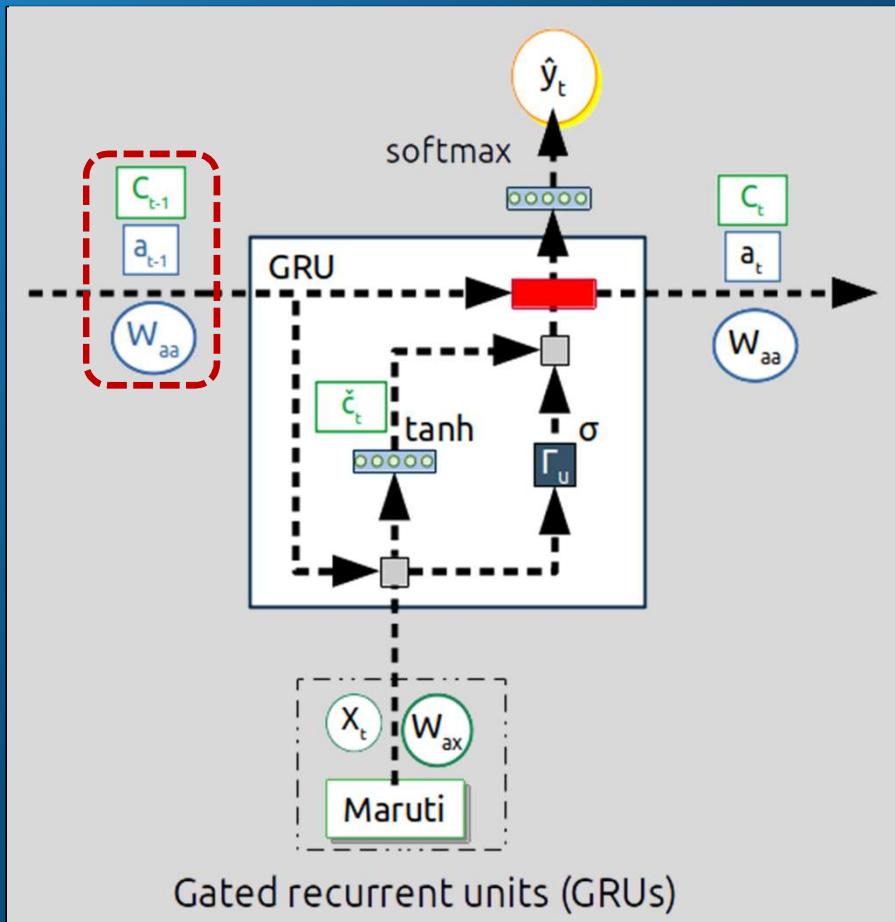


□ RNN Equation

- ❖ $a_t = g_1 ([a_{t-1} : x_t] \cdot W_a + b_a)$
and
- ❖ $\hat{y}_t = g_2 (a_t \cdot W_y + b_y)$

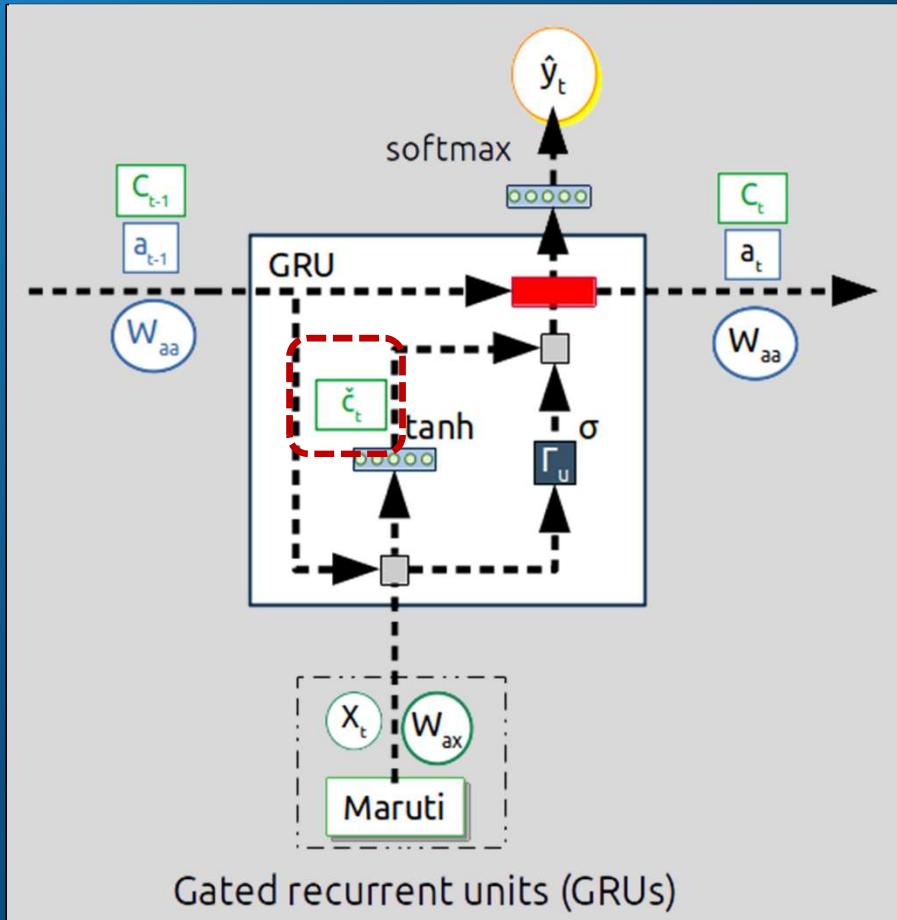
□ Alternate layout for one of the recurrence

GRU Cell



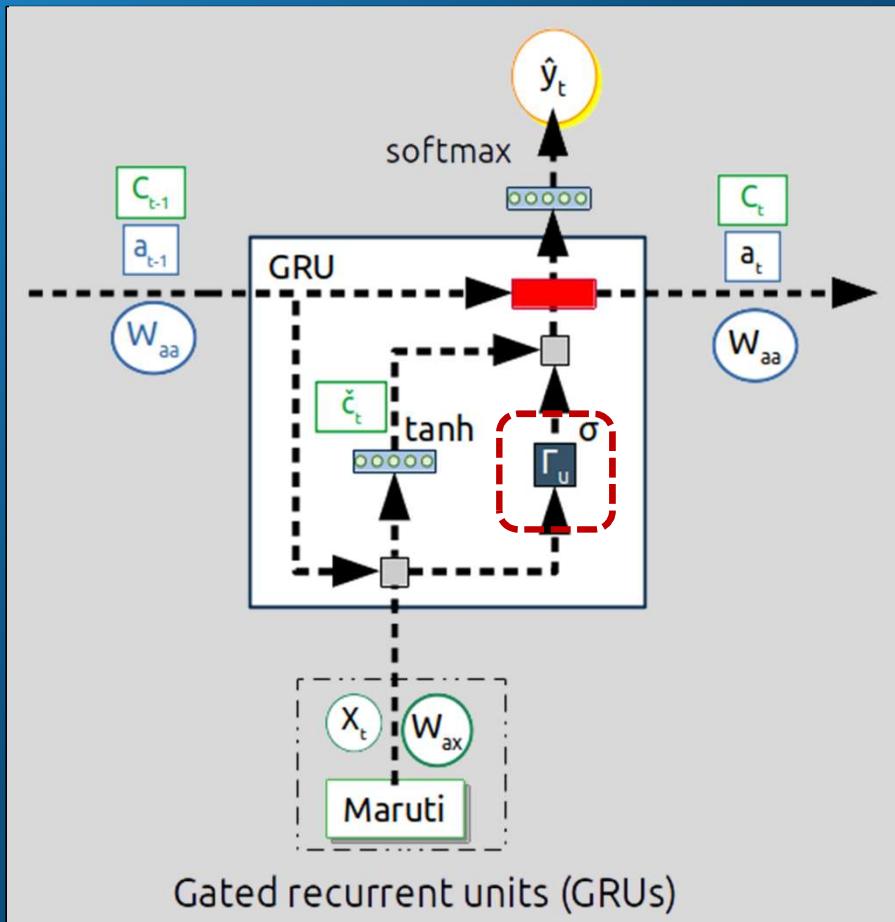
- GRU will have a cell which we will use to **retain** values from earlier iterations.
- Cell is called memory cell and is represented by 'c'.
- In this case , memory cell value 'c' and activation 'a' value will be same.

GRU Cell



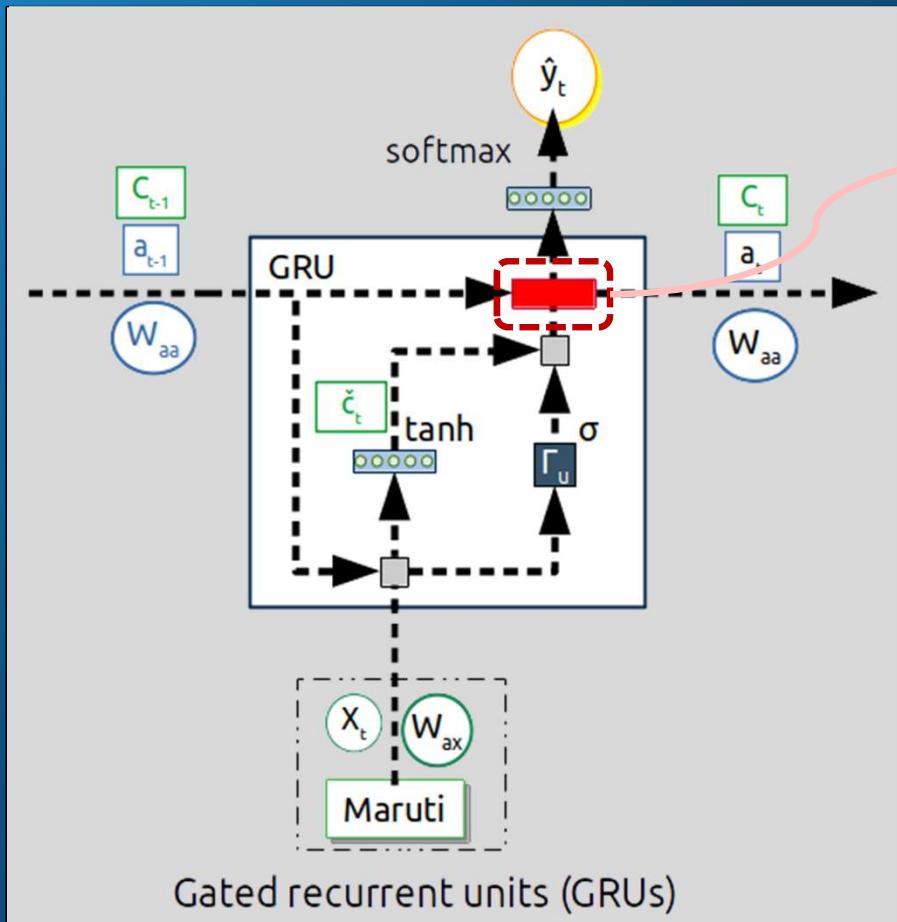
- At every time step, we will calculate ' \hat{c} ', a candidate which may replace ' c_{t-1} '
❖ $\hat{c}_t = \tanh ([c_{t-1} : x_t] \cdot W_c + b_c)$
- May be not...
- Next we need some parameter which will tell us if we need to replace the value of ' c_t ' with candidate ' \hat{c}_t ' or not
- Note: At present both c_t and a_t are same... keeping them separate for consistency

GRU Cell



- Important change is called ' Gate'
- Gate, represented by 'Γ' Gamma is:
 - ❖ $\Gamma_u = \sigma ([c_{t-1} : x_t] \cdot W_u + b_u)$
- In above equation, we are using separate Weights and Biases for update gate.
- The activation used here is sigmoid so for most part our Gamma will be 1 or 0

GRU Cell

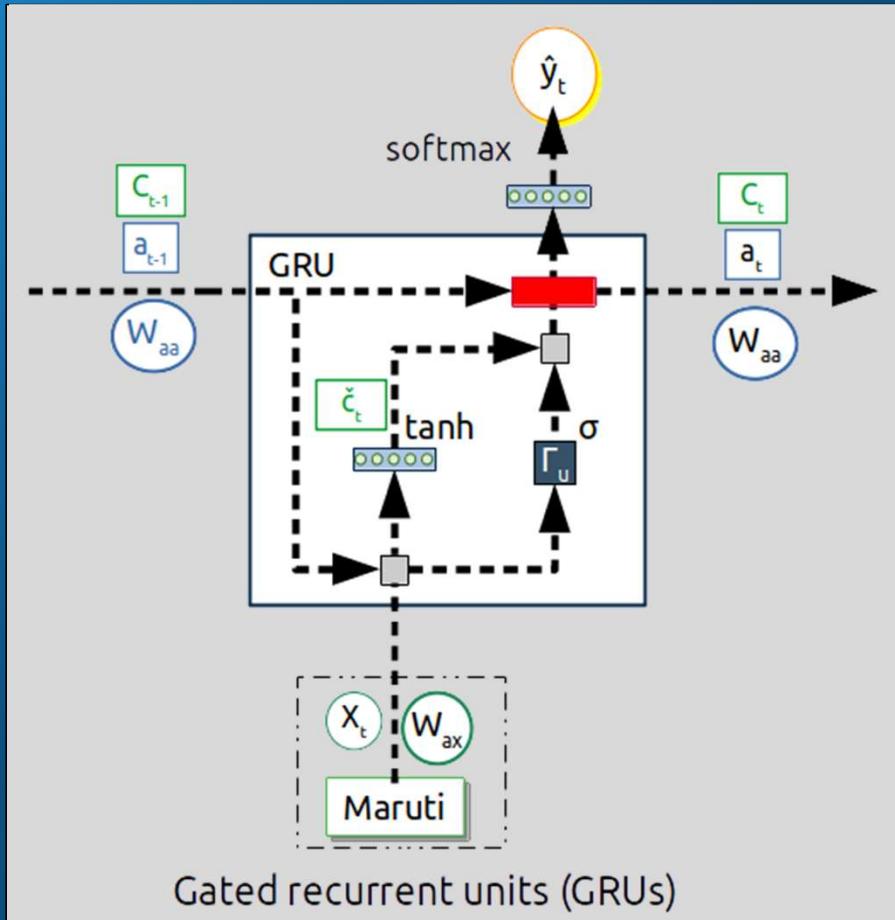


- With ' Γ_u ' Gamma update, we can calculate c_t as follows:
 - ❖ $c_t = \Gamma_u * \hat{c}_t + (1-\Gamma_u) * c_{t-1}$
 - ❖ The Red box represents the above equation

- Intuition: how Γ_u will be used????
- “I **felt happy** because I saw the others **were happy**”
- and because I knew I should **feel happy**, but I **wasn’t really happy**.”

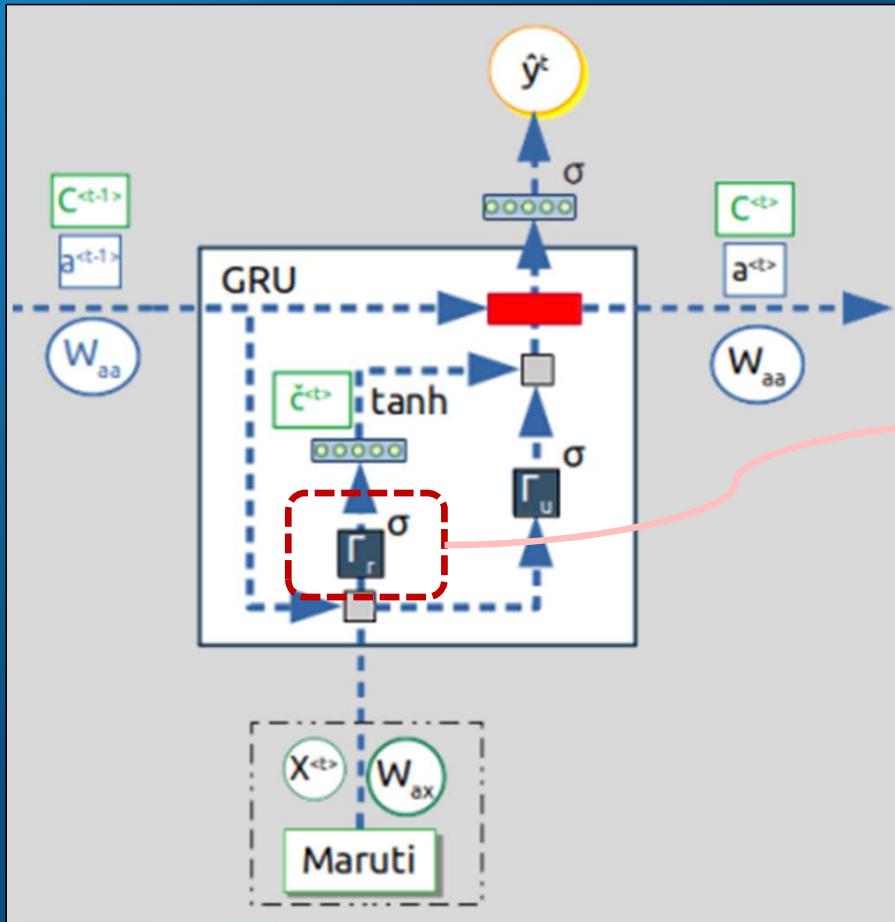
- Create or keep c_t
- Time to Replace c_t

GRU Cell



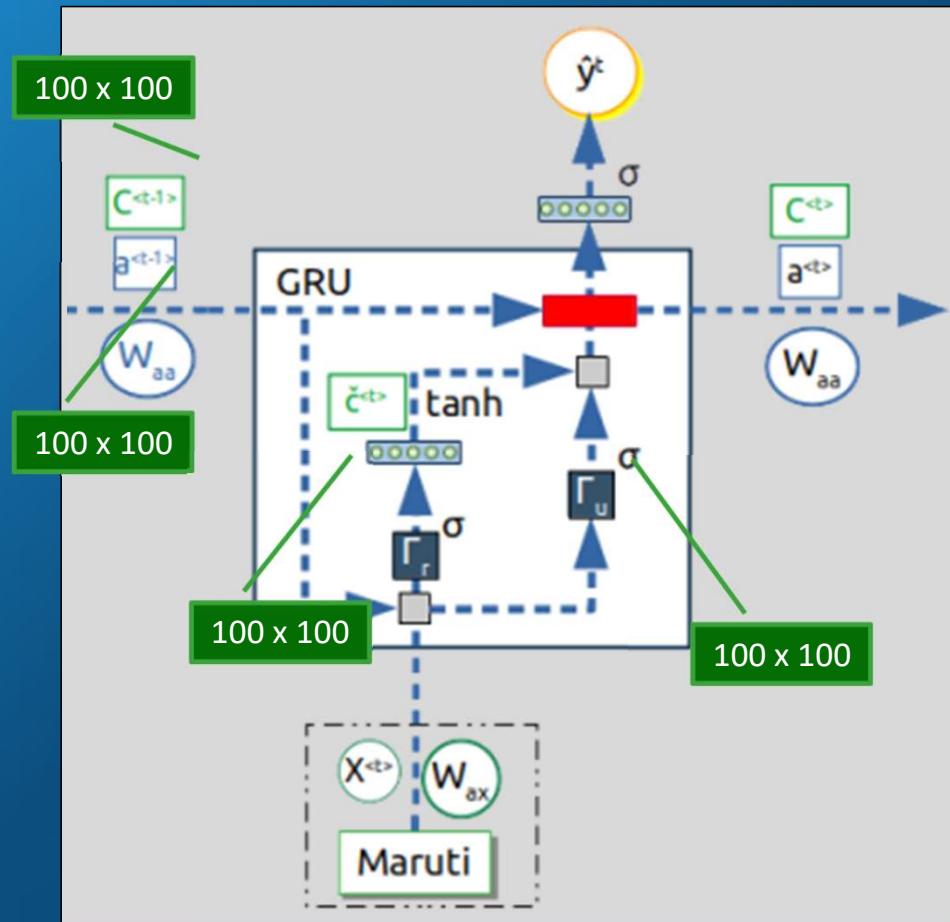
- ❑ So we can maintain c_t and their values through many layers and use them repeatedly.
 - ❑ And more importantly, when to change it
 - ❑ I **felt happy** because I saw the others **were happy** and because I knew I should **feel happy**, but I **wasn't really happy**.

GRU Cell



- That was GRU in simplified form.
- Actual implementation has one additional parameter Γ_r Gamma Relevance
 - ❖ $\hat{c}_t = \tanh([c_{t-1}: x_t]. W_c + b_c)$
 - ❖ $\Gamma_u = \sigma([c_{t-1}: x_t]. W_u + b_u)$
 - ❖ $\Gamma_r = \sigma([c_{t-1}: x_t]. W_r + b_r)$
 - ❖ $c_t = \Gamma_u * \hat{c}_t + (1 - \Gamma_u) * c_{t-1}$

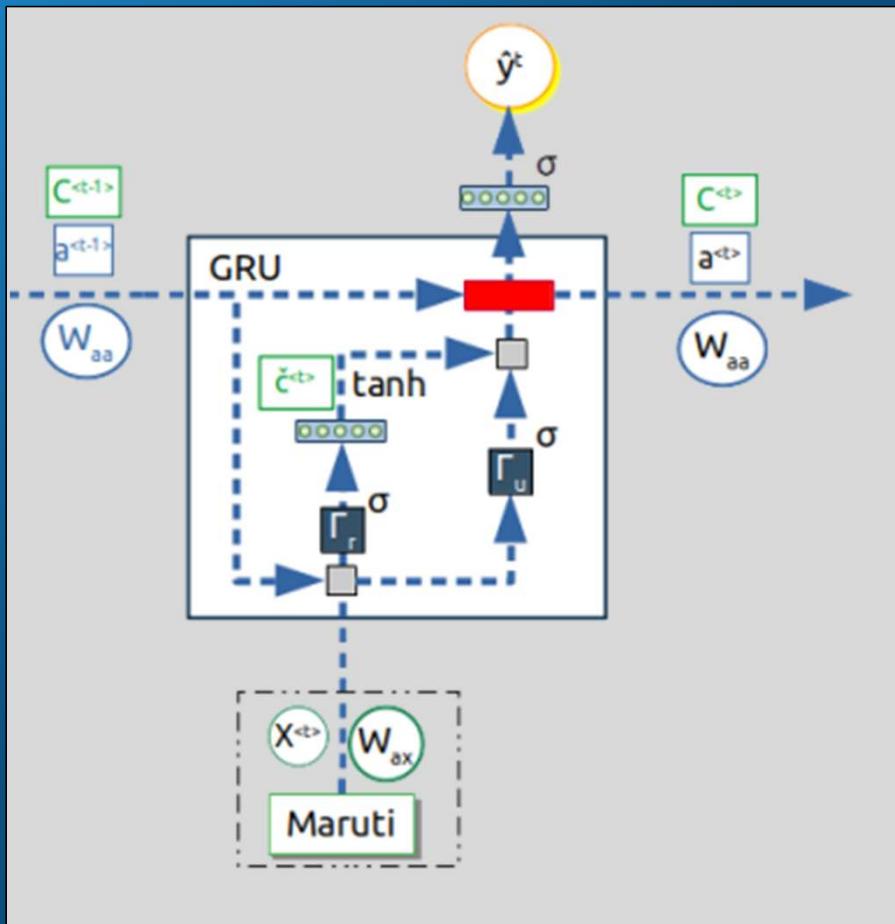
GRU Cell



- That was GRU in simplified form.
- Actual implementation has one additional parameter Γ_r Gamma Relevance
 - ❖ $\hat{c}_t = \tanh ([\Gamma_r * c_{t-1} : x_t] . W_c + b_c)$
 - ❖ $\Gamma_u = \sigma ([c_{t-1} : x_t] . W_u + b_u)$
 - ❖ $\Gamma_r = \sigma ([c_{t-1} : x_t] . W_r + b_r)$
 - ❖ $c_t = \Gamma_u * \hat{c}_t + (1 - \Gamma_u) * c_{t-1}$

Element wise multiplications

GRU Cell



Extended GRU:

$$\check{c}_t = \tanh([\Gamma_r * c_{t-1} : x_t] \cdot W_c + b_c)$$

$$\Gamma_u = \sigma([c_{t-1} : x_t] \cdot W_u + b_u)$$

$$\Gamma_r = \sigma([c_{t-1} : x_t] \cdot W_r + b_r)$$

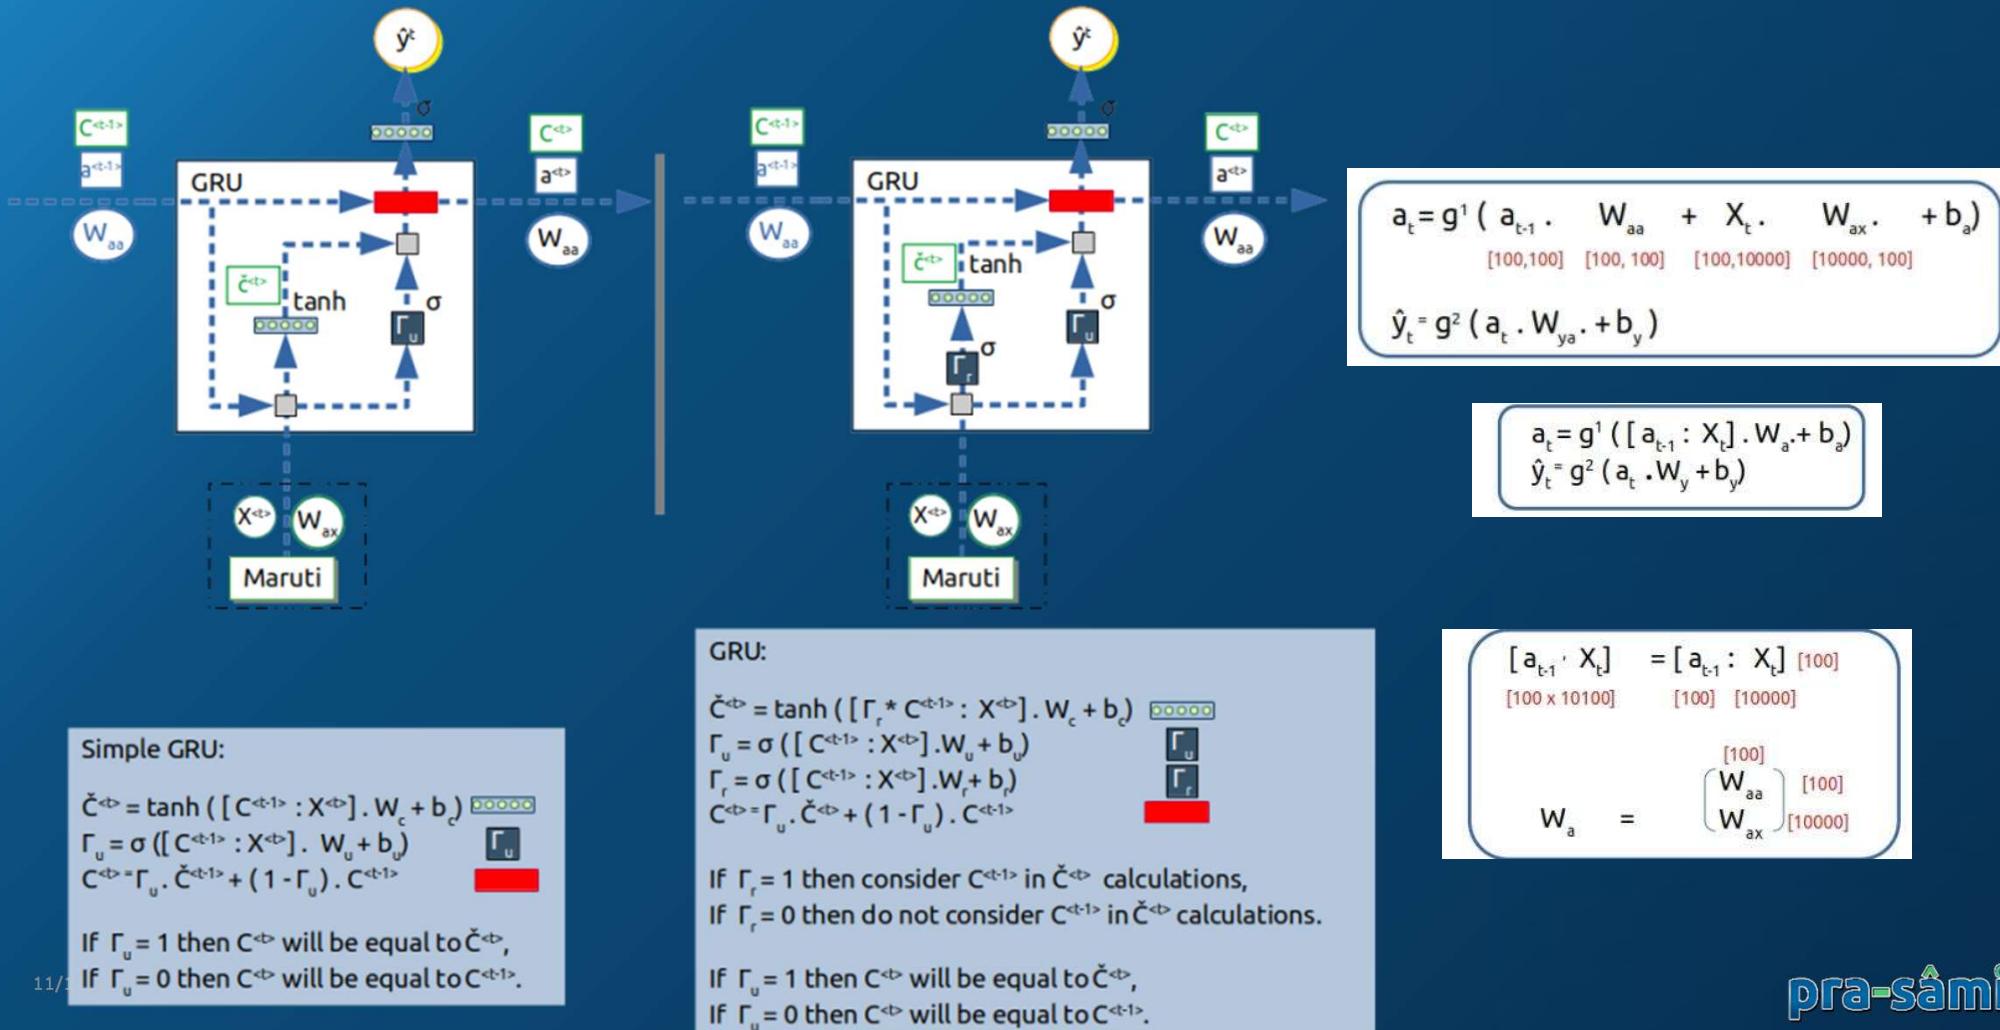
$$c_t = \Gamma_u \cdot \check{c}_t + (1 - \Gamma_u) \cdot c_{t-1}$$



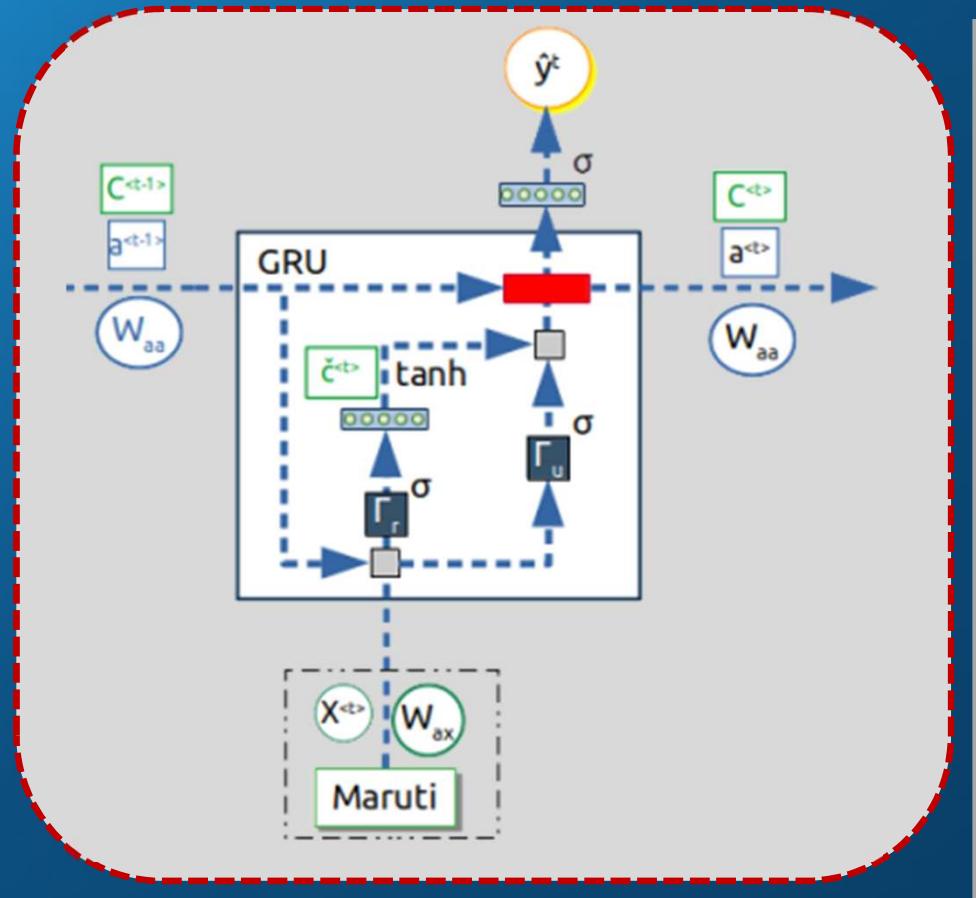
If $\Gamma_u = 1$ then c_t will be equal to \check{c}_t ,

If $\Gamma_u = 0$ then c_t will be equal to c_{t-1}

GRU Cell



GRU Cell



□ And that, my friends, is GRU....

$$\hat{c}_t = \tanh ([\Gamma_r * c_{t-1} : x_t] \cdot W_c + b_c)$$

$$\Gamma_u = \sigma ([c_{t-1} : x_t] \cdot W_u + b_u)$$

$$\Gamma_r = \sigma ([c_{t-1} : x_t] \cdot W_r + b_r)$$

$$c_t = \Gamma_u * \hat{c}_t + (1 - \Gamma_u) * c_{t-1}$$

□ Coming up next... LSTM!

